



**Leibniz Supercomputing Centre**  
of the Bavarian Academy of Sciences and Humanities

# AI Training Series Introduction to the LRZ AI Systems

03.05.2023 | Maja Piskac

# ANNOUNCEMENT



<https://app1.edoobox.com/en/LRZ/>

Registration deadline is today at noon!

| Offer Name   | Date                    | Type                  | Price  | Places | Action                       |
|--|-------------------------|-----------------------|--------|--------|------------------------------|
| Introduction to ANSYS CFX on LRZ HPC Systems                                     | 16.03.2023 – 27.04.2023 | ONLINE                | 0.00 € | 62     | <a href="#">Register now</a> |
| AI Training Series - Orientation Session   | 12.04.2023 – 13.04.2023 | HYBRID: ONLINE/LRZ    | 0.00 € | 30     | <a href="#">Register now</a> |
| EuroCC AI for Science Bootcamp (apply via openhackathons.org)                    | 17.04.2023 – 18.04.2023 | ONLINE                | 0.00 € | 100    | <a href="#">Register now</a> |
| Introduction to LRZ HPC Systems with Focus on CFD Workflows                      | 19.04.2023 – 19.04.2023 | ONLINE                | 0.00 € | 51     | <a href="#">Register now</a> |
| AI Training Series - Intro to Container Technology & Application to AI at LRZ    | 26.04.2023 – 26.04.2023 | HYBRID: ONLINE/LRZ    | 0.00 € | 43     | <a href="#">Register now</a> |
| AI Training Series - Introduction to the LRZ AI Systems                          | 03.05.2023 – 03.05.2023 | ONLINE                | 0.00 € | 47     | <a href="#">Register now</a> |
| AI Training Series - Intel AI Workshop - Accelerated Machine Learning with Intel | 04.05.2023 – 04.05.2023 | ONLINE                | 0.00 € | 88     | <a href="#">Register now</a> |
| SuperMUC-NG Status and Results Workshop  | 09.05.2023 – 11.05.2023 | ONLINE                | 0.00 € | 192    | <a href="#">Register now</a> |
| AI Training Series - Introduction to the LRZ Linux Cluster                       | 09.05.2023 – 09.05.2023 | HYBRID: ONLINE/LRZ    | 0.00 € | 73     | <a href="#">Register now</a> |
| Working with Noisy Quantum Computers   | 11.05.2023 – 11.05.2023 | Leibniz Rechenzentrum | 0.00 € | 3      | <a href="#">Register now</a> |
| AI Training Series - High Performance Data Analytics Using R at LRZ              | 12.05.2023 – 12.05.2023 | HYBRID: ONLINE/LRZ    | 0.00 € | 83     | <a href="#">Register now</a> |
| EuroCC N-Ways to GPU Programming Bootcamp (apply via openhackathons.org)         | 15.05.2023 – 16.05.2023 | ONLINE                | 0.00 € | 100    | <a href="#">Register now</a> |
| AI Training Series - Intel AI Workshop - Accelerated Deep Learning with Intel    | 16.05.2023 – 16.05.2023 | HYBRID: ONLINE/LRZ    | 0.00 € | 95     | <a href="#">Register now</a> |
| AI Training Series - Introduction to the LRZ Compute Cloud                       | 17.05.2023 – 17.05.2023 | HYBRID: ONLINE/LRZ    | 0.00 € | 85     | <a href="#">Register now</a> |

## 1. Introduction to the LRZ AI Systems

---

- Overview of the LRZ AI Systems
- Access to the LRZ AI Systems
- NVIDIA NGC Cloud
- Introduction to Enroot Containers
- Interactive and Batch Jobs
- Open on Demand
- Exercise: Run a job with an Enroot container

## 2. Fundamentals of Deep Learning

---

- Introduction to Neural Networks
- Introduction to Convolutional Neural Networks
- Exercises: Train CNNs on a GPU
- Introduction to Transformers
- Exercise: Train a Transformer on a GPU

## 3. Distributed Training of Neural Networks

---

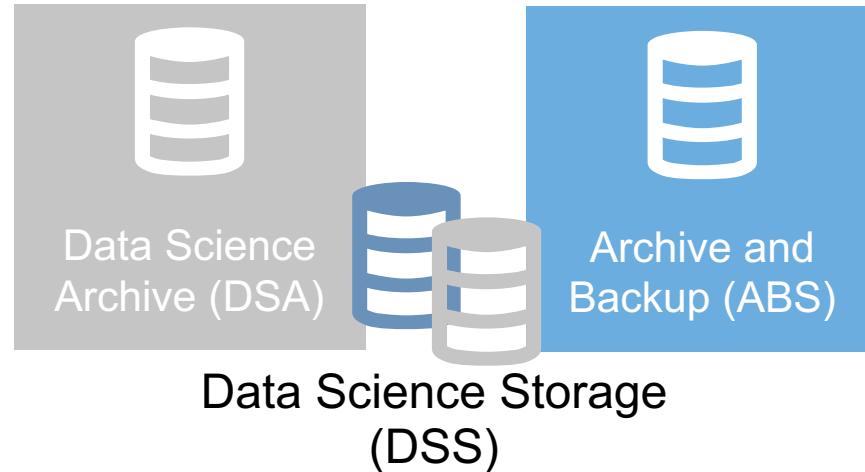
- Data Parallel Training
- Exercise: DP Training of CNN on 2 GPUs
- Model Parallel Training: Pipeline Parallel and Tensor Parallel
- Exercise: PP Training of Transformer on 2 GPUs

## Planned breaks

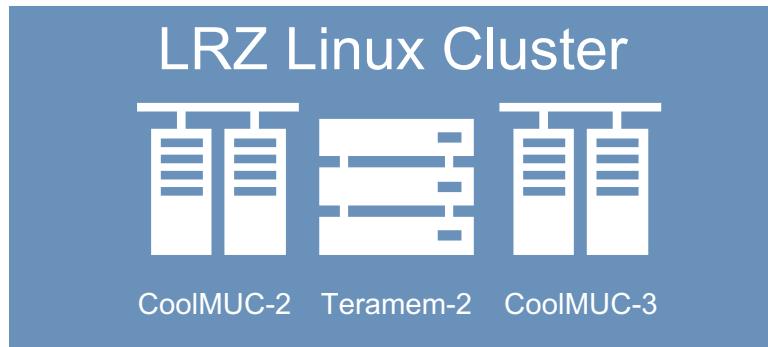
- 11:15 – 11:30 Coffee Break I
- 12:30 – 13:30 Lunch break
- 14:45 – 15:00 Coffee Break II

# 1. Introduction to the LRZ AI Resources

## Overview of the LRZ Systems

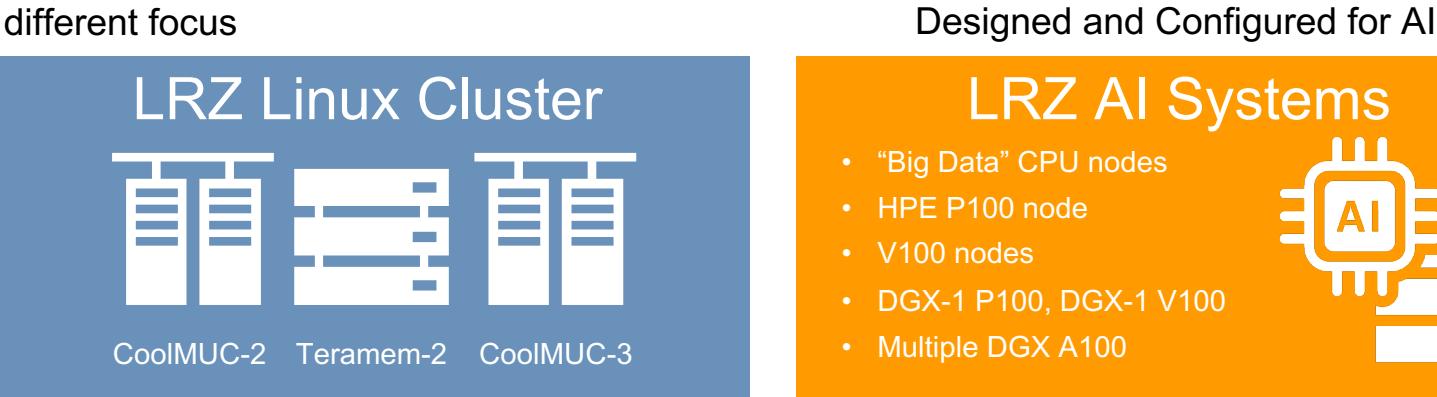


Multi-purpose cluster systems might be used for AI workloads as well, but have different focus



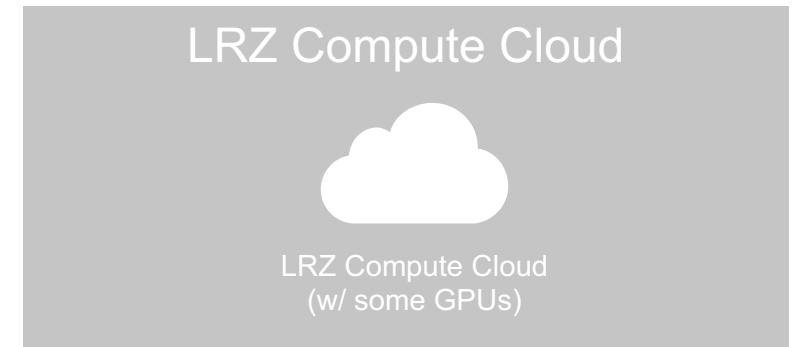
`lxlogin[1-3].lrz.de`

`lxlogin8.lrz.de`



`login.ai.lrz.de`  
`https://datalab3.srv.lrz.de`

Flexible system that copes with almost any workload



`https://cc.lrz.de`

# 1. Introduction to the LRZ AI Resources

## Overview of the LRZ AI Systems



|                           | DGX A100 Architecture          | DGX A100 Architecture MIG         | DGX-1 V100 Architecture | DGX-1 P100 Architecture | HPE Intel Skylake + Nvidia Node | V100 GPU Nodes      | CPU Nodes     |
|---------------------------|--------------------------------|-----------------------------------|-------------------------|-------------------------|---------------------------------|---------------------|---------------|
| <b>Number of Nodes</b>    | 4                              | 1                                 | 1                       | 1                       | 1                               | 4                   | 9             |
| <b>CPU cores per node</b> | 256                            | 256                               | 80                      | 80                      | 64                              | 20                  | 40 / 32 / 20  |
| <b>Memory per node</b>    | 2 TB                           | 1 TB                              | 512 GB                  | 512 GB                  | 256 GB                          | 368 GB              | min. 360 GB   |
| <b>GPUs per node</b>      | 8 NVIDIA A100                  | 8 NVIDIA A100 (16 MIG partitions) | 8 Nvidia Tesla V100     | 8 Nvidia Tesla P100     | 4 Nvidia Tesla P100             | 2 Nvidia Tesla V100 | --            |
| <b>Memory per GPU</b>     | 80 GB                          | 40 GB (20GB)                      | 16 GB                   | 16 GB                   | 16GB                            | 16 GB               | --            |
| <b>SLURM Partition</b>    | lrz-dgx-a100-80x8              | lrz-dgx-a100-40x8                 | lrz-dgx-1-v100x8        | lrz-dgx-1-p100x8        | lrz-hpe-p100x4                  | lrz-v100x2          | lrz-cpu       |
| <b>Nodes</b>              | lrz-dgx-a100-[001-002,004-005] | lrz-dgx-a100-003                  | dgx-002                 | dgx-001                 | p100-001                        | gpu-[001-003,005]   | cpu-[001-009] |

# 1. Introduction to the LRZ AI Resources

## DGX A100



### 1. 8x NVIDIA A100 GPUs with up to 640GB total GPU memory

12 NVIDIA NVLinks® per GPU, 600GB/s of GPU-to-GPU bidirectional bandwidth

### 2. 6x NVIDIA NVSwitches™

4.8TB/s of bidirectional bandwidth, 2X more than previous-generation NVSwitch

### 3. 10x NVIDIA ConnectX-7 200Gb/s network interface

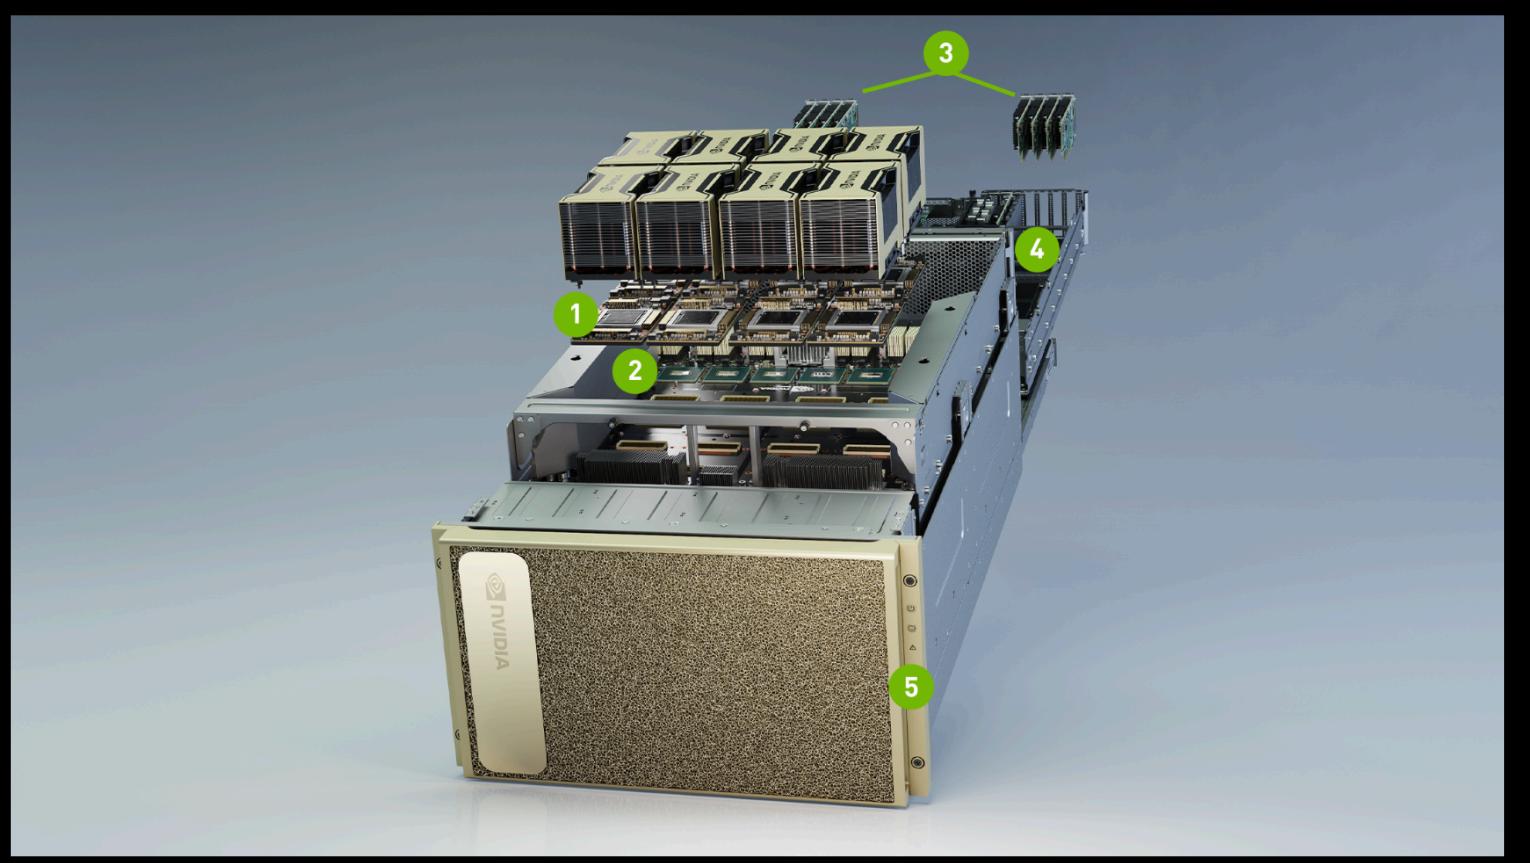
500GB/s of peak bidirectional bandwidth

### 4. Dual 64-CORE AMD CPUs and 2TB system memory

3.2X more cores to power the most intensive AI jobs

### 5. 30TB Gen4 NVMe SSD

50GB/s of peak bandwidth, 2X faster than Gen3 NVMe SSDs



## 1. Introduction to the LRZ AI Resources

# Storage on the LRZ AI Systems



| Storage Pool                         | Designated Use   | Top-level Directory                         | Size Limit                    | Automated Backup                              | Expiration                    | Additional Information                                       |
|--------------------------------------|--|---|-------------------------------|---|-------------------------------|--|
| <b>Home directory</b>                | infrequently changing user files, scripts, configuration files, etc.<br>unified home directory with the LRZ Linux Cluster, created when LRZ Linux Cluster access is granted<br><b>! Not suitable for heavy and/or high-frequency I/O operations, use the AI Systems DSS instead.</b> | /dss/dsshome1/.../<br><user>                | 100 GB                        | Yes, backup to tape and file system snapshots | Lifetime of LRZ project       | <a href="#">File Systems and IO on Linux-Cluster</a>         |
| <b>AI Systems DSS</b>                | high-bandwidth, low latency I/O, access is granted upon request through the LRZ Servicedesk  | /dss/dssfs04                                | up to 5 TB                    | No  | Until further notice          |  |
| <b>Linux Cluster DSS</b>             | general purpose, long-term data storage  | /dss/dssfs02<br>/dss/dssfs03                | up to 10 TB                   | No  | Lifetime of data project      | <a href="#">File Systems and IO on Linux-Cluster</a>         |
| <b>Exclusive/private DSS systems</b> | specified by the system owner, can be purchased, implemented and housed exclusively, for a private group of dedicated users  | /dsslegfs01<br>/dsslegfs02<br>/dssmcmrlfs01 | specified by the system owner | specified by the system owner                 | specified by the system owner | <a href="#">Data Science Storage</a> ("joint project offer") |

# Access to the LRZ AI Systems – How to access?

- User requirements to get the access:
  1. Own / get a Linux Cluster account:  
<https://doku.lrz.de/display/PUBLIC/Access+and+Login+to+the+Linux-Cluster>
  2. Submitt a service request to [LRZ Servicedesk](#) – select "AI topics" and "LRZ AI Systems - Request for Access" from the drop-down lists. Request has to include a description of the intended usage.

- Login node [login.ai.lrz.de](https://login.ai.lrz.de) accessible via ssh:

```
$ ssh --login_name=xxyyzz login.ai.lrz.de
```

- Resources are allocated and jobs submitted to the partitions using a job scheduler Slurm.

<https://slurm.schedmd.com/overview.html>

## 1. Introduction to the LRZ AI Resources

# Access to the LRZ AI Systems – sinfo, salloc, srun & scancel



```
$ ssh --login_name=xxyyzzz login.ai.lrz.de
```

Executes in the login node login-1

```
$ sinfo
```

```
$ squeue --user=xxyyzzz
```

```
$ salloc --partition=lrz-v100x2 --gres=gpu:1
```

```
$ srun --pty bash
```

```
$ scancel job_id
```

```
di82hod@login-1:~$ sinfo
PARTITION          AVAIL  TIMELIMIT  NODES  STATE NODELIST
lrz-v100x2*        up    14-00:00:0  1      mix  gpu-005
lrz-v100x2*        up    14-00:00:0  3      alloc  gpu-[001-003]
lrz-hpe-p100x4     up    14-00:00:0  1      idle  p100-001
lrz-dgx-1-p100x8   up    14-00:00:0  1      drain* dgx-001
lrz-dgx-1-v100x8   up    14-00:00:0  1      alloc  dgx-002
lrz-dgx-a100-80x8  up    14-00:00:0  2      mix  lrz-dgx-a100-[002,004]
lrz-dgx-a100-80x8  up    14-00:00:0  2      alloc  lrz-dgx-a100-[001,005]
lrz-dgx-a100-40x8-mig up    14-00:00:0  1      mix  lrz-dgx-a100-003
lrz-cpu            up    14-00:00:0  6      mix  cpu-[001-002,004-005,007,009]
lrz-cpu            up    14-00:00:0  1      alloc  cpu-003
lrz-cpu            up    14-00:00:0  2      idle  cpu-[006,008]
mcml-dgx-a100-40x8 up    14-00:00:0  2      mix  mcml-dgx-[004,008]
mcml-dgx-a100-40x8 up    14-00:00:0  6      alloc  mcml-dgx-[001-003,005-007]
test-v100x2         up    14-00:00:0  1      idle  gpu-004
di82hod@login-1:~$ salloc -p lrz-hpe-p100x4 --gres=gpu:2 --time=02:00:00!
salloc: Pending job allocation 162333
salloc: job 162333 queued and waiting for resources
salloc: job 162333 has been allocated resources
salloc: Granted job allocation 162333
di82hod@login-1:~$ srun --pty bash
di82hod@p100-001:~$ exit
exit
di82hod@login-1:~$ scancel 162333
di82hod@login-1:~$ salloc: Job allocation 162333 has been revoked.

di82hod@login-1:~$
```

# 1. Introduction to the LRZ AI Resources

## Nvidia NGC Containers



- The NGC catalog provides access to GPU-accelerated software that speeds up end-to-end workflows with performance-optimized containers, pretrained AI models, and SDKs that can be deployed on any NVIDIA's GPU-powered on-prem, cloud and edge systems.

- <https://catalog.ngc.nvidia.com>

The screenshot shows the NVIDIA NGC Catalog homepage. At the top, there is a search bar labeled "Search NGC Catalog...". Below the search bar, there are four main sections under the heading "Getting Started":

- HPC Collection**: Collection - High Performance Computing. Description: This collection provides access to the top HPC applications for Molecular Dynamics, Quantum Chemistry, and Scientific visualization. Includes a "View Labels" button.
- NVIDIA AI - End-to-End AI Development ...**: Collection - Deep Learning. Description: This is a collection of performance-optimized frameworks, SDKs, and models to build Computer Vision and Speech AI applications. Includes a "View Labels" button.
- NGC - Getting Started**: Collection - Beginner. Description: Looking to get started with containers and models on NGC? This is the place to start. Includes a "View Labels" button.
- Deep Learning Frameworks**: Collection - Deep Learning. Description: This collection contains performance-optimized AI frameworks including PyTorch and TensorFlow. Includes a "View Labels" button.

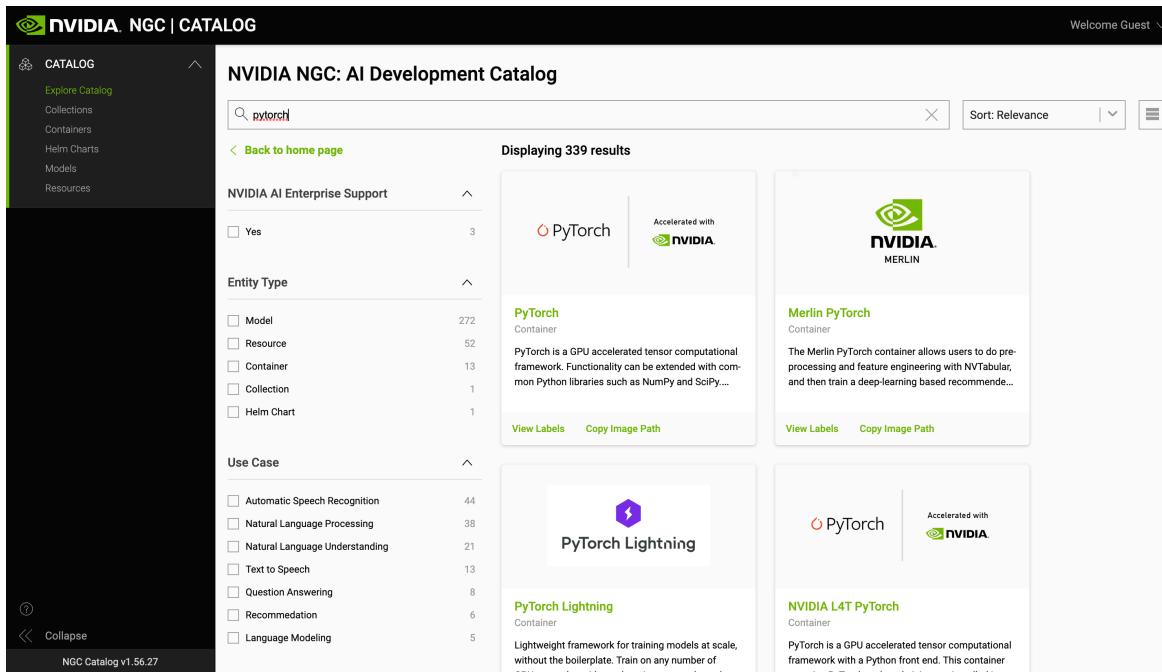
Below these sections, there are three additional cards:

- Documentation**: Description: We've got a whole host of documentation, covering the NGC UI and our powerful CLI. You can find out more here. Includes a "Go to Documentation" button.
- Command Line Interface**: Description: Want to get more from NGC? Everything you see here can be used and managed via our powerful CLI tools. Includes a "Download Now" button.
- NGC Private Registry**: Description: Private Registries from NGC allow you to secure, manage, and deploy your own assets to accelerate your journey to AI. Includes a "Learn More" button.

At the bottom, there is a section titled "Popular Collections" with a "See All Collections" link.

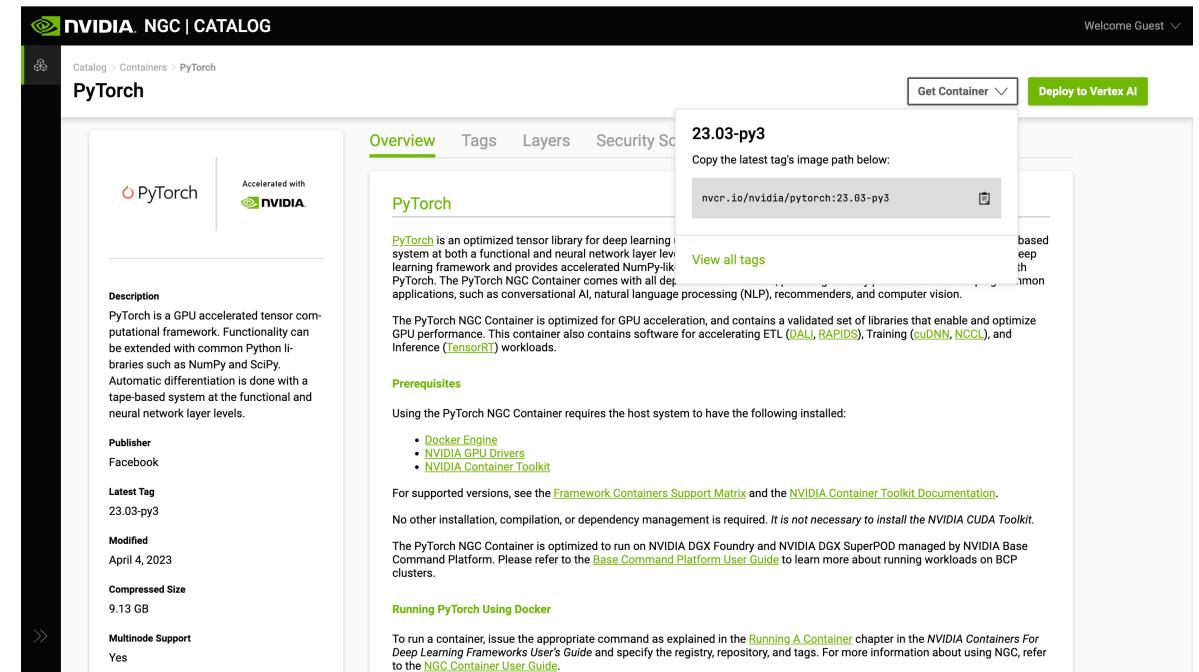
# 1. Introduction to the LRZ AI Resources

## Nvidia NGC Containers



The screenshot shows the NVIDIA NGC Catalog interface. A search bar at the top contains the query "pytorch". Below it, a message says "Displaying 339 results". On the left, there are filters for "Entity Type" (Model, Resource, Container, Collection, Helm Chart) and "Use Case" (Automatic Speech Recognition, Natural Language Processing, Natural Language Understanding, Text to Speech, Question Answering, Recommendation, Language Modeling). The main area lists four container options:

- PyTorch Container**: Accelerated with NVIDIA. Description: PyTorch is a GPU accelerated tensor computational framework. Functionality can be extended with common Python libraries such as NumPy and SciPy....
- Merlin PyTorch Container**: Accelerated with NVIDIA. Description: The Merlin PyTorch container allows users to do pre-processing and feature engineering with NVTabular, and then train a deep-learning based recommende...
- PyTorch Lightning Container**: Accelerated with NVIDIA. Description: Lightweight framework for training models at scale, without the boilerplate. Train on any number of GPUs in parallel without changing code and tr...
- NVIDIA L4T PyTorch Container**: Accelerated with NVIDIA. Description: PyTorch is a GPU accelerated tensor computational framework with a Python front end. This container contains PyTorch and torchvision are installed in a...



The screenshot shows the details page for the PyTorch container. At the top, there are buttons for "Get Container" and "Deploy to Vertex AI". The main content includes:

- Overview**: PyTorch is an optimized tensor library for deep learning, system at both a functional and neural network layer level, learning framework and provides accelerated NumPy-like PyTorch. The PyTorch NGC Container comes with all default applications, such as conversational AI, natural language processing (NLP), recommenders, and computer vision.
- Tags**: 23.03-py3
- Layers**: Copy the latest tag's image path below: `nvr.io/nvidia/pytorch:23.03-py3`
- Security Score**: Based on the latest tag, the score is 100%.
- Description**: PyTorch is a GPU accelerated tensor computational framework. Functionality can be extended with common Python libraries such as NumPy and SciPy. Automatic differentiation is done with a tape-based system at the functional and neural network layer levels.
- Publisher**: Facebook
- Latest Tag**: 23.03-py3
- Modified**: April 4, 2023
- Compressed Size**: 9.13 GB
- Multinode Support**: Yes
- Running PyTorch Using Docker**: To run a container, issue the appropriate command as explained in the [Running A Container](#) chapter in the [NVIDIA Containers For Deep Learning Frameworks User's Guide](#) and specify the registry, repository, and tags. For more information about using NGC, refer to the [NGC Container User Guide](#).

# 1. Introduction to the LRZ AI Resources

## Nvidia NGC Containers – Setting up credentials



The screenshot shows the NVIDIA NGC Catalog homepage. At the top right, there is a user profile for "Welcome Guest". Below it, there are links for "Setup", "Terms of Use", and "Sign In/Sign Up". A large green button labeled "Register for NGC" is prominently displayed. The main content area features a large image of a neural network and text about deploying performance-optimized AI/HPC software containers. Below this, there is a section titled "NVIDIA NGC: AI Development Catalog" with a message "Displaying 0 results" and a "Sorting & Filters" button. Under "Getting Started", there are two small thumbnail images.

The screenshot shows the "Setup" section of the NVIDIA NGC Setup page. It features a "Generate API Key" section with a "Get API Key" button and a "CLI" section with "Documentation" and "Downloads" buttons. The top right corner shows a user profile for "Maja Piskac".

The screenshot shows the NVIDIA NGC Catalog homepage, similar to the one above but with a different layout. It features the "NVIDIA LaunchPad" section with a "Get Started" button and a message "Instantly experience end-to-end workflows with free hands-on labs". Below this is the "NVIDIA NGC: AI Development Catalog" section with "Displaying 0 results" and "Sorting & Filters" buttons. The "Getting Started" section includes two thumbnail images.

The screenshot shows the "API Key" page under "Setup > API Key". It has sections for "API Information" (describing how it authenticates NGC service), "Usage" (instructions for logging in with an API key), and "NGC CLI" (examples of commands like \$ ngc config set and \$ docker login nvcr.io). The top right corner shows a user profile for "Maja Piskac".

### Nvidia NGC Containers – Setting up credentials



- Create the file `enroot/.credentials` within your `$HOME` and append the following line to it:

```
machine nvcr.io login $oauthtoken password <Your Key>
```

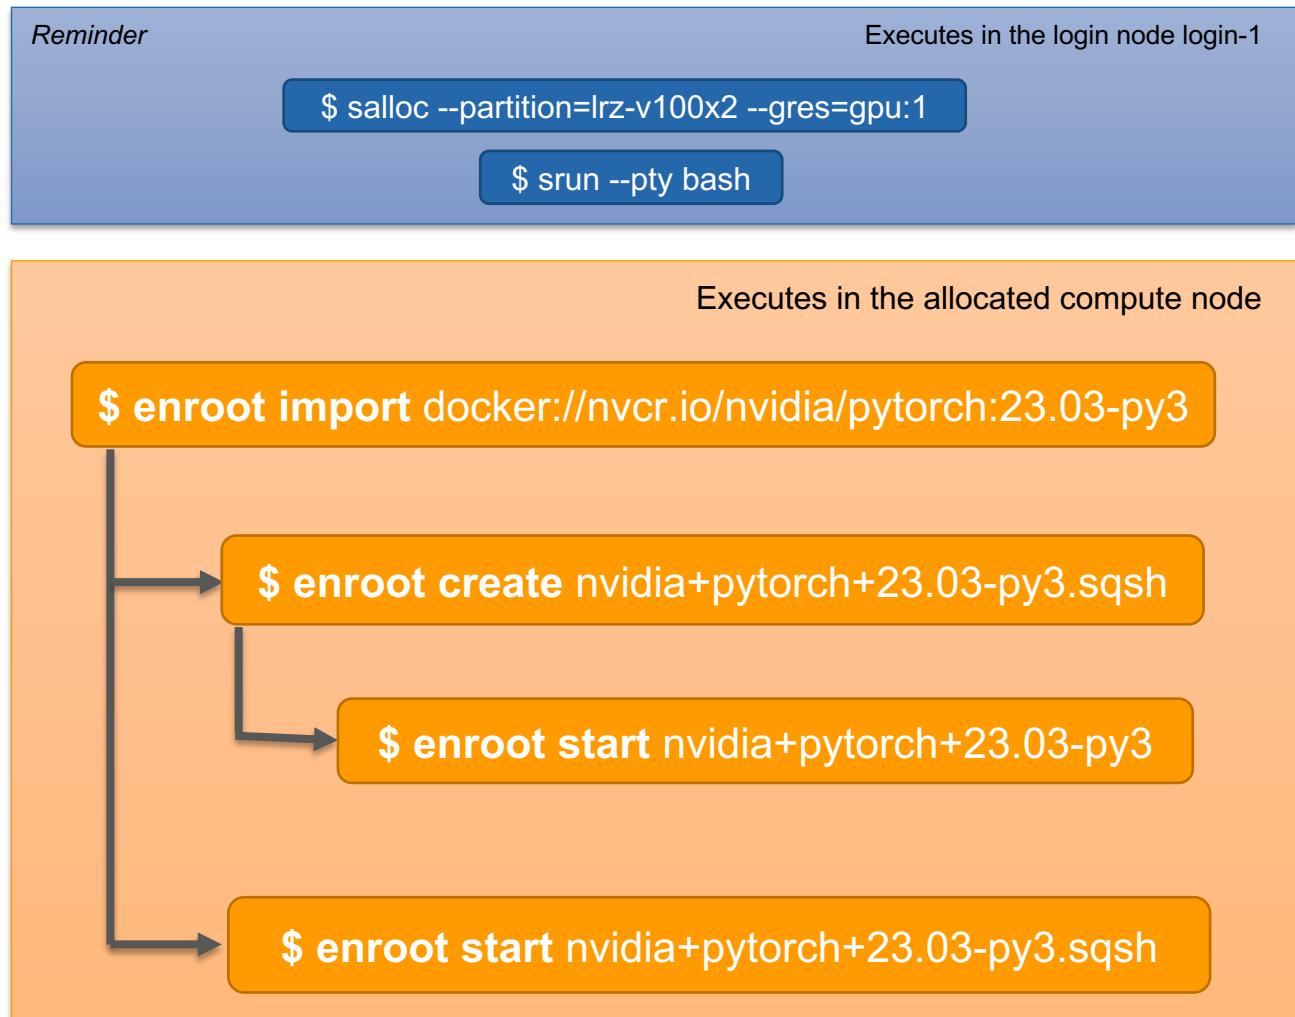
- Introduce a new line after `<Your Key>`
- Now you can import containers from Nvidia NGC, e.g. latest Pytorch container, with the following line:

```
$ enroot import docker://nvcr.io/nvidia/pytorch:23.03-py3
```

- These containers supply the cuda toolkit, cudnn libraries, and Nvidia dependencies.

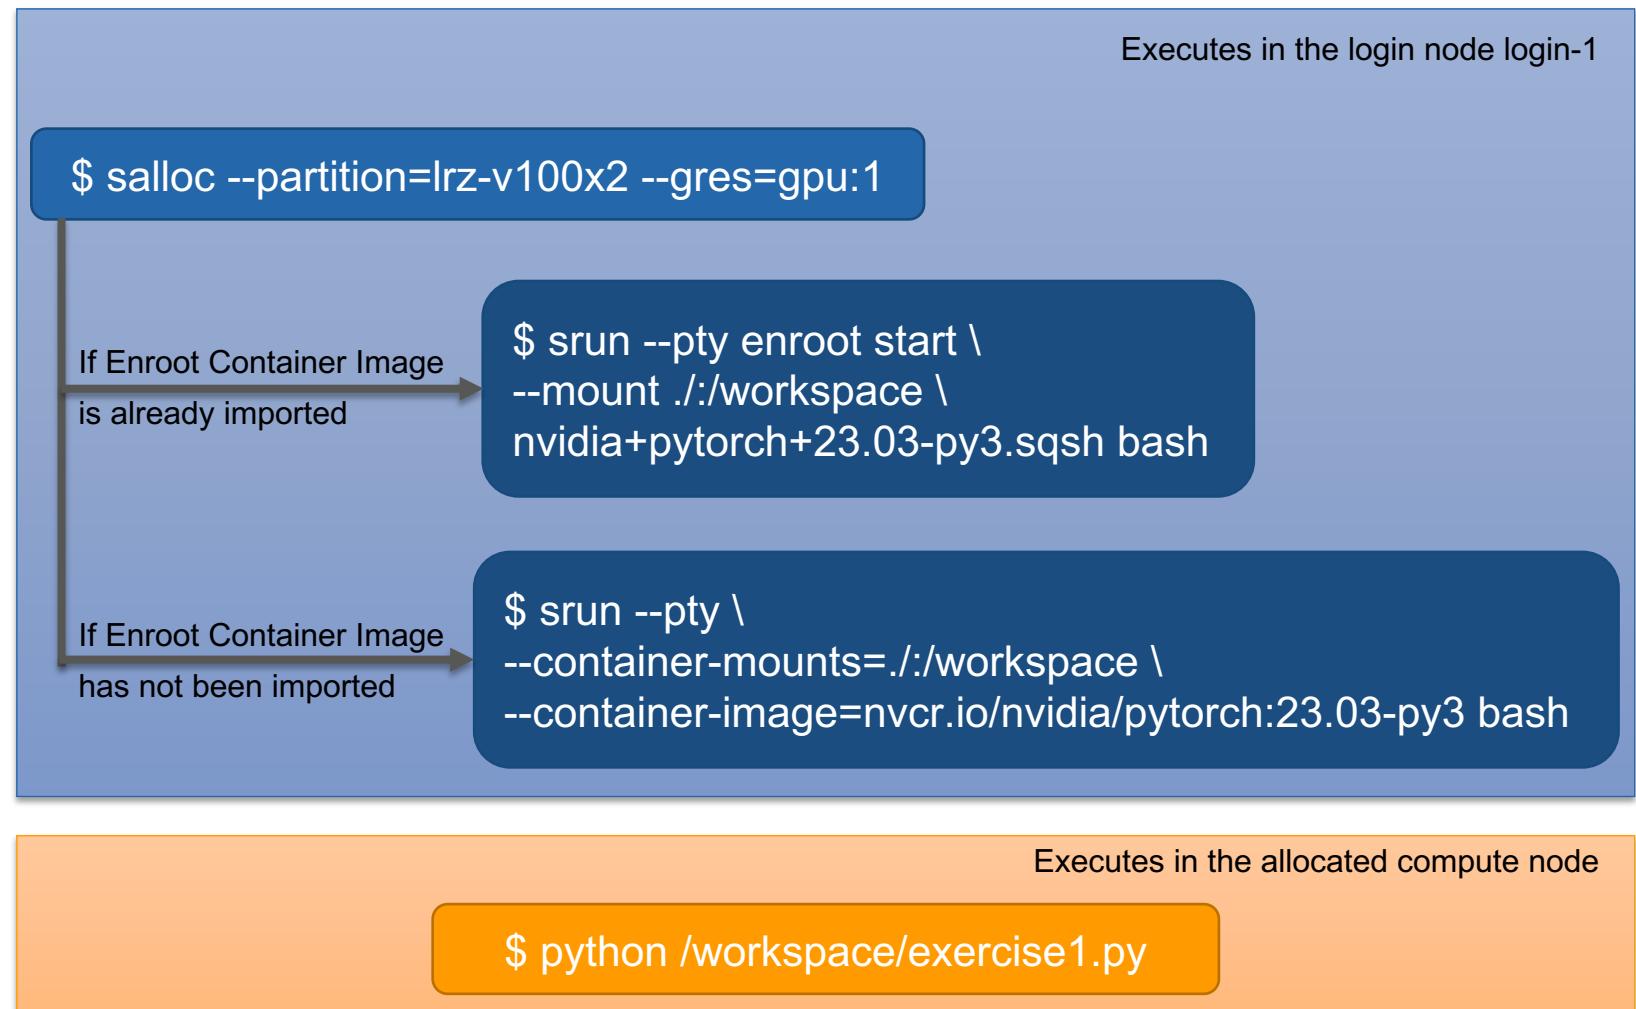


- Enroot container runtime operates completely in user space.
  - It allows to run containers defined by container images from the NVIDIA NGC Cloud or from the Docker Hub.
  - Not available on the login node, but on the compute nodes!
  - The Enroot Workflow:
    - I. **Import** an Enroot Container Image – resulting in **sqsh** file
    - II. Create an Enroot Container with **create**,
    - III. Run software inside an existing Enroot Container with **start**.



# Running Applications as Interactive Jobs

- Interactive jobs are submitted to an existing allocation of resources using the **srun** command.
- We can **mount** existing data from outside of the container into the container.
- Enroot container creation and job submission in a single step can be done via a plugin called *pyxis*.



# Running Applications as Batch Jobs

- Batch jobs are the preferred and quicker way of using the LRZ AI Systems.
- Batch job is queued and executed when the resources are available.
- It does the allocation and running of the job for you (instead of salloc and srun).
- The **sbatch** command submits jobs described in a **sbatch script file**.
- You need to specify the partition or nodes that you want to use.
- Two additional required arguments: *output* and *error messages* file.

```
#!/bin/bash
#SBATCH -p lrz-dgx-a100-80x8
#SBATCH --gres=gpu:1
#SBATCH -o exercise1.out
#SBATCH -e exercise1.err

srun \
--container-mounts='.:/:workspace' \
--container-image='nvcr.io/nvidia/pytorch:23.03-py3' \
python /workspace/exercise1.py
```

Executes in the login node login-1

\$ sbatch exercise1.sbatch

# Dealing with base images from catalogues other than NGC

- If your chosen image does not supply the cuda toolkit, do not install it within the image!  
Let the Enroot runtime deal with it by adding environment variables within the container.

Executes in the allocated resource

```
$ enroot import docker://someone/image-no-cuda
```

```
$ enroot create --name base_container someone+image-no-cuda.sqsh
```

```
$ enroot start base_container bash
```

```
echo "NVIDIA_DRIVER_CAPABILITIES=compute,utility,video" >> /etc/environment  
echo "NVIDIA_REQUIRE_CUDA=cuda>=9.0" >> /etc/environment  
echo "NVIDIA_VISIBLE_DEVICES=all" >> /etc/environment
```

```
$ exit
```

```
$ enroot export --output base_image.sqsh base_container
```

### Creating an extended Enroot image

- If your workload depends on a package not provided by the used image.

Executes in the allocated resource

```
$ enroot create --name plt_container base_image.sqsh
```

```
$ enroot start plt_container
```

```
$ pip install matplotlib
```

```
$ exit
```

```
$ enroot export --output pytorch+plt.sqsh plt_container
```

```
$ enroot start pytorch+plt.sqsh
```

# Access to AI Systems through interactive web servers



- Jupyter Notebook, JupyterLab, RStudio Server and TensorBoard
- Available at <https://datalab3.srv.lrz.de>
- For a typical use-case:
  - select the type of resources (CPU only or CPU + single GPU)
  - specify your workload (a combination of CPU core and RAM requirements)
  - select the container environment you want to work with
  - finally, specify the number of hours you plan to work (be aware that your session will be shut down when this time limit is reached, any unsaved work will then be lost).

# 1. Introduction to the LRZ AI Resources

## Access to AI Systems through interactive web servers



The screenshot illustrates the LRZ AI OnDemand interface, which provides integrated access to various AI resources. The main window shows a navigation sidebar with options like 'Sh', 'An', 'Anm', 'LRZ', 'Ben', 'Pass', and 'OnDemand'. The 'OnDemand' section contains a link to 'rv.lrz.de/pun/sys/dashboard#'. A central panel displays the 'Systems Web UI (TEST INSTANCE)' with tabs for 'Files', 'Jobs', 'Clusters', 'Interactive Apps', and a 'Servers' dropdown. The 'Interactive Apps' tab is selected, showing options for 'Jupyter Notebook', 'RStudio Server', and 'TensorBoard'. A sub-modal window titled 'Jupyter Notebook' is open, showing a Jupyter Notebook interface with a URL 'datalab3.srv.lrz.de/node/p100-001.srv.lrz.de:8942/notebooks/ondemand/Exercise1.py.ipynb'. The notebook content includes Python code for importing packages, defining hyper-parameters, and setting up a neural network architecture.

```
# Import packages
import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms

# Hyper-parameters
image_width = 32
image_channels = 3
conv1_out_channels = 50
conv2_out_channels = 75
kernel_size = 5
pool_size = 2
fc1_out_channels = 50
num_classes = 10
num_epochs = 5
batch_size = 100
learning_rate = 0.001
dim_1 = int((image_width-kernel_size+1)/pool_size)
dim_2 = int((dim_1-kernel_size+1)/pool_size)
```

## 1. Introduction to the LRZ AI Resources

# Public Datasets and Containers on the LRZ AI Systems



- LRZ AI Systems offer a dedicated Data Science Storage (DSS) container aimed at storing public datasets as well as Enroot container images of interest to more than one researcher.

| Dataset       | Location  | Version  | Licence   |
|---------------|---|--|---|
| AlphaFold     | /dss/dssfs04/pn69za/pn69za-dss-0004/datasets/alphafold/   | Last update April 2022 following the instructions here <a href="https://github.com/deepmind/alphafold#genetic-databases">https://github.com/deepmind/alphafold#genetic-databases</a> | <a href="https://github.com/deepmind/alphafold#license-and-disclaimer">https://github.com/deepmind/alphafold#license-and-disclaimer</a> |
| COCO-Stuff    | dss/dssfs04/pn69za/pn69za-dss-0004/datasets/cocostuff/    | 2020 Update (train2017.zip, val2017.zip, annoations_trainval2017.zip, stuff_annotations_trainval2017.zip)  | <a href="https://cocodataset.org/#termsofuse">https://cocodataset.org/#termsofuse</a>   |
| Visual Genome | dss/dssfs04/pn69za/pn69za-dss-0004/datasets/visualgenome/ | Version 1.4 of dataset completed as of July 12, 2017   | Creative Commons Attribution 4.0 International License  |

| Container | Location   | Version   |
|-----------|--|---|
| AlphaFold | /dss/dssfs04/pn69za/pn69za-dss-0004/containers/shared/alphafold.sqsh | Built from the Dockerfile here<br><a href="https://github.com/deepmind/alphafold/tree/main/docker">https://github.com/deepmind/alphafold/tree/main/docker</a> |

# 1. Introduction to the LRZ AI Resources

## Summary

### Interactive job

Command line

```
$ ssh -l xxxyyzz login.ai.lrz.de
```

Executes in the login node login-1

```
$ sinfo
```

```
$ salloc --partition=dgx-1-p100 --gres=gpu:1
```

```
$ srun --pty --container-mounts='./workspace' \
--container-image=nvcr.io/nvidia/pytorch:23.03-py3 bash
```

Executes in the allocated resource

```
$ python /workspace/exercise1.py
```

### Batch job

```
#!/bin/bash
#SBATCH -p lrz-dgx-a100-80x8
#SBATCH --gres=gpu:1
#SBATCH -o exercise1.out
#SBATCH -e exercise1.err
```

```
srun --container-mounts='./workspace' \
--container-image='nvcr.io/nvidia/pytorch:23.03-py3' \
python /workspace/exercise1.py
```

Executes in the login node datalab2

```
$ sbatch exercise1.sbatch
```

Web browser

<https://datalab3.srv.lrz.de>

### Exercise: Start an interactive job and create an Enroot container



1. Ssh into the login node [login.ai.lrz.de](https://login.ai.lrz.de) and download the exercises from <https://doku.lrz.de/display/PUBLIC/AI+Infrastructure+Material>.
2. Allocate and run a job with 1 GPU.
3. Import a Pytorch container image from NGC (or Pytorch Docker container), create and start the container.
4. Run exercise1.py as an interactive job.

## 1. Introduction to the LRZ AI Systems

---

- ✓ Overview of the LRZ AI Systems
- ✓ Access to the LRZ AI Systems
- ✓ NVIDIA NGC Cloud
- ✓ Introduction to Enroot Containers
- ✓ Interactive and Batch Jobs
- ✓ Open on Demand
- ✓ Exercise: Run a job with an Enroot container

## 2. Fundamentals of Deep Learning

---

- Introduction to Neural Networks
- Introduction to Convolutional Neural Networks
- Exercises: Train CNNs on a GPU
- Introduction to Transformers
- Exercise: Train a Transformer on a GPU

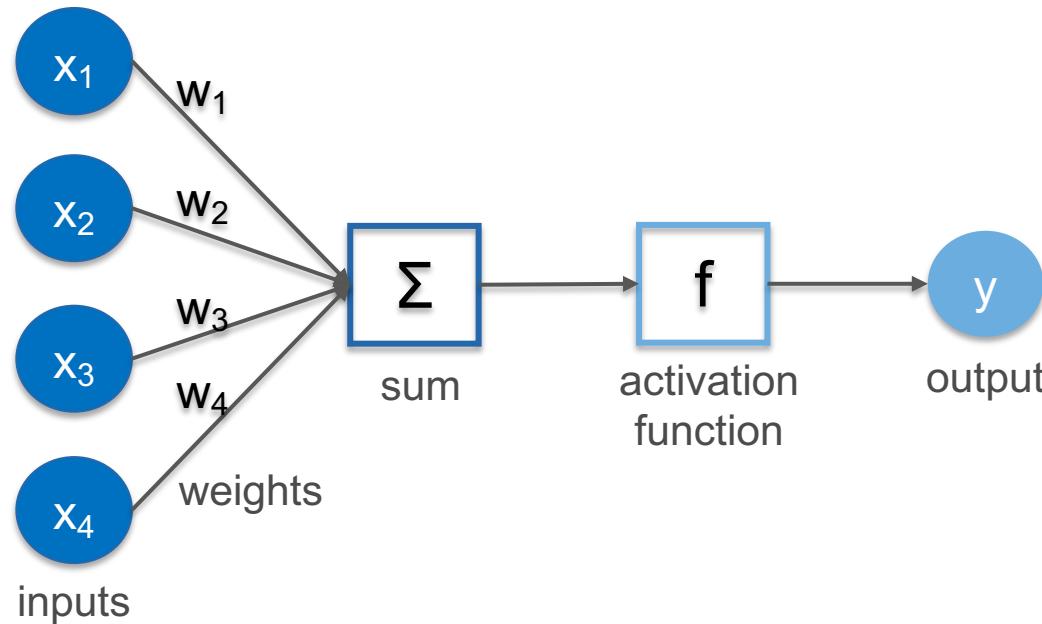
## 3. Distributed Training of Neural Networks

---

- Data Parallel Training
- Exercise: DP Training of CNN on 2 GPUs
- Model Parallel Training: Pipeline Parallel and Tensor Parallel
- Exercise: PP Training of Transformer on 2 GPUs

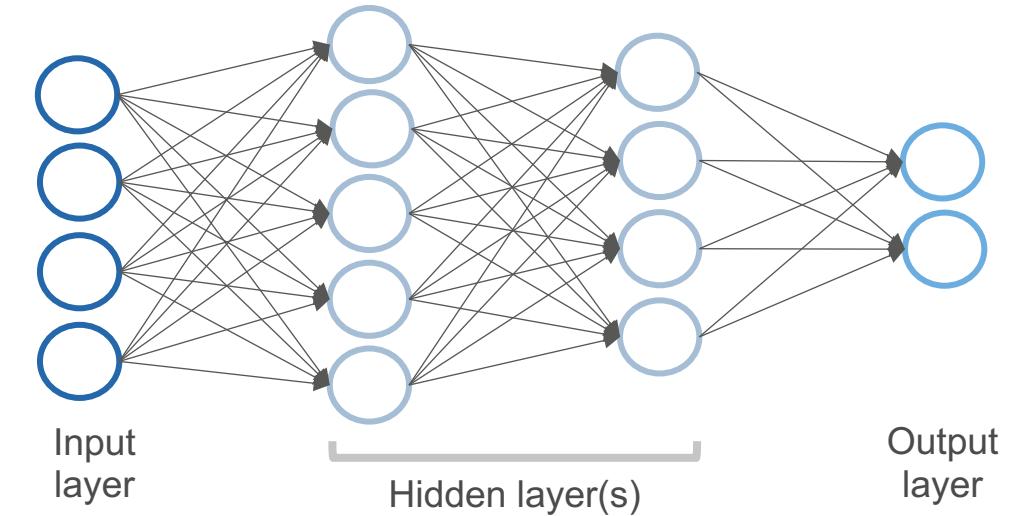
## 2. Fundamentals of Deep Learning

# Introduction to Neural Networks



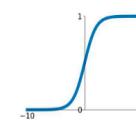
$$y = f(x_1 * w_1 + x_2 * w_2 + x_3 * w_3 + x_4 * w_4 + b)$$

$$\theta = \{w_1, w_2, w_3, w_4\}$$



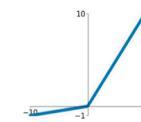
Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



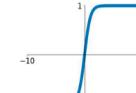
Leaky ReLU

$$\max(0.1x, x)$$



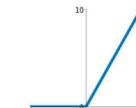
tanh

$$\tanh(x)$$



ReLU

$$\max(0, x)$$

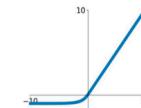


Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



## 2. Fundamentals of Deep Learning

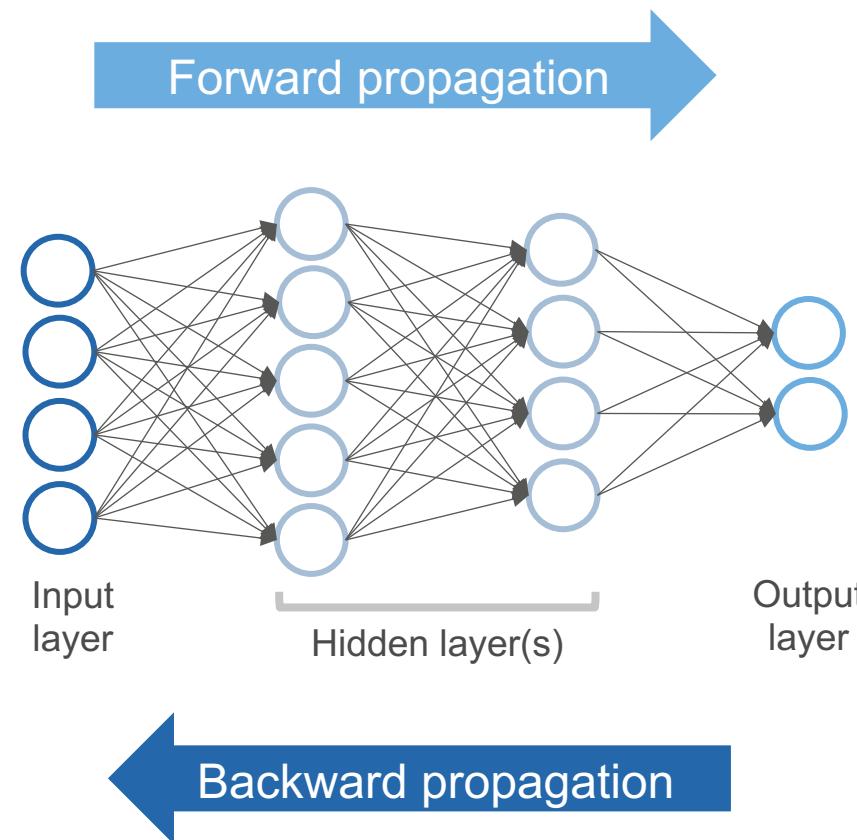
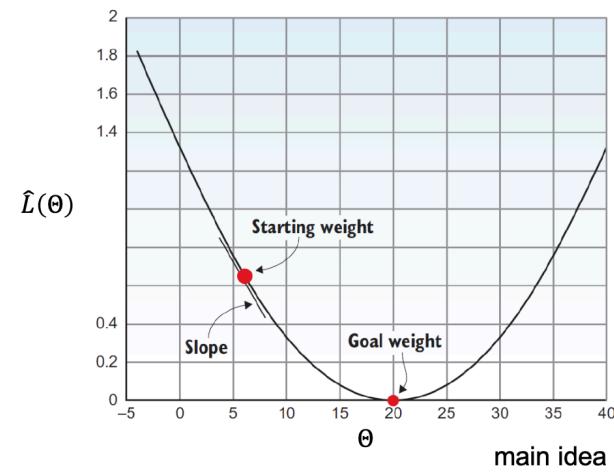
### Training Neural Networks

- Find weights that minimize the loss function on a given subset of the training data.

#### Optimizer SGD

$$\theta_{new} = \theta_{old} - \eta * g(\theta_{old})$$

$$g(\theta) = \frac{1}{n} \sum_{i=1}^n \nabla l(\theta; z_i)$$



#### Loss function

$$L(\theta) = \frac{1}{n} \sum_{i=1}^n l(\theta; z_i)$$

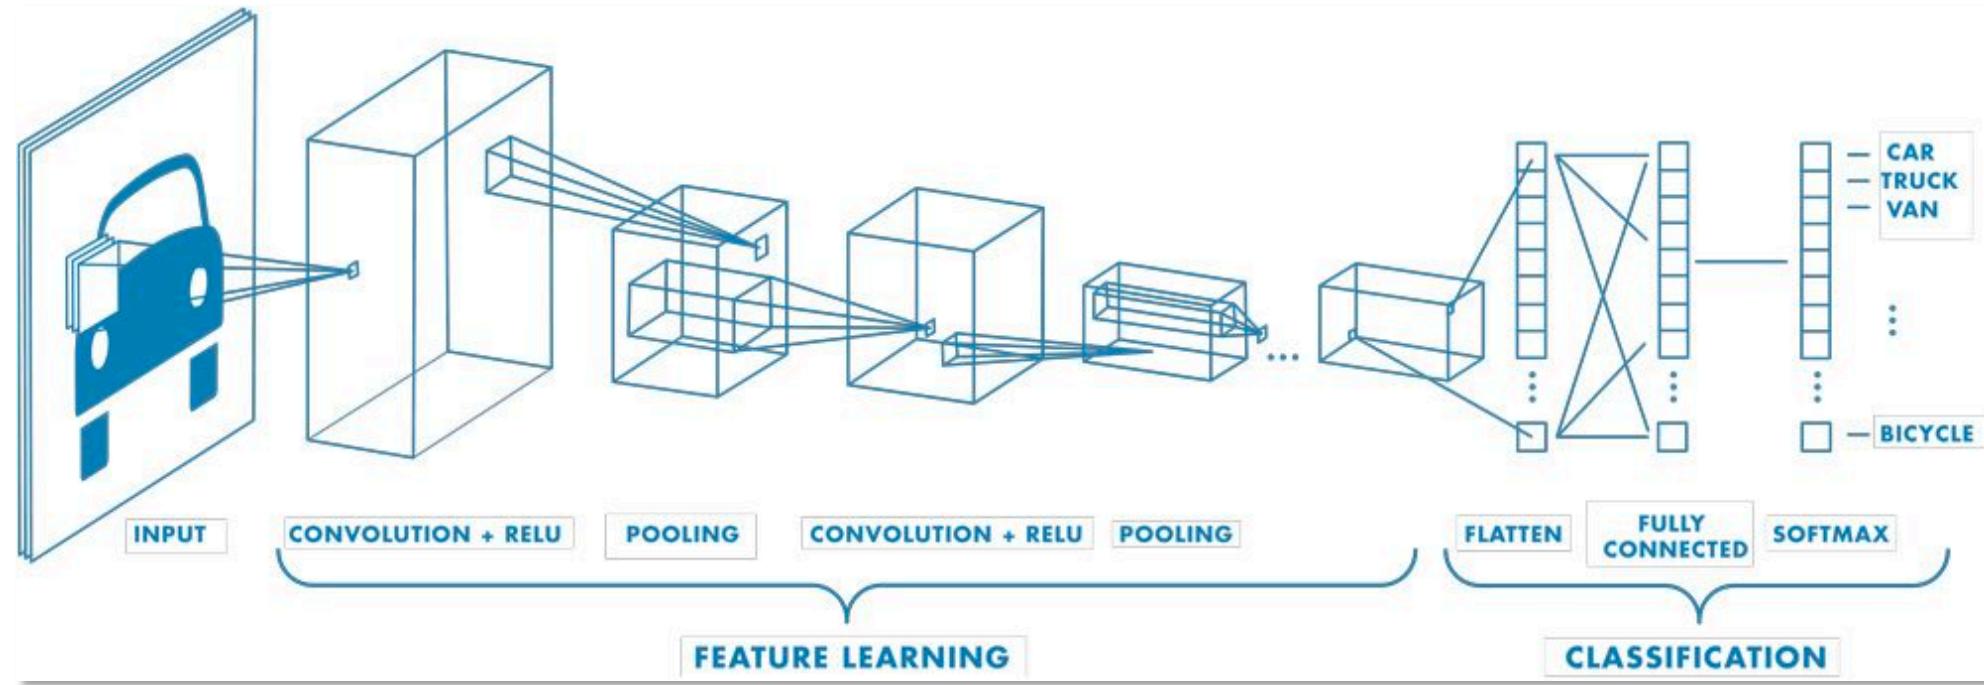
$$MeanSquaredError(y, y^p) = \frac{1}{n} \sum_{i=1}^n (y_i - y_i^p)^2$$

$$CrossEntropy(y, y^p) = -\frac{1}{n} \sum_{i=1}^n \sum_{c=1}^M y_{i,c} * \log(y_{i,c}^p)$$

#### Hyper Parameters

- epochs
- batch size  $n$
- learning rate  $\eta$

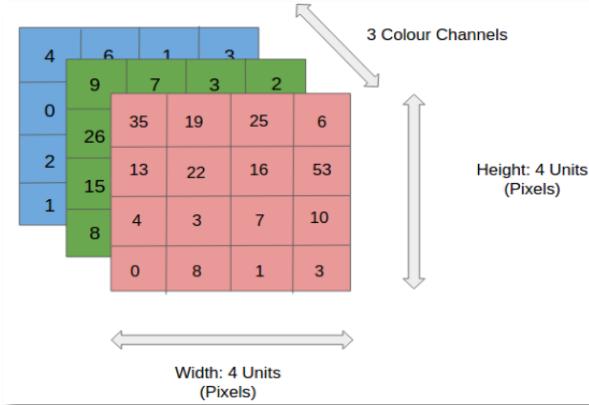
# Introduction to Convolutional Neural Networks



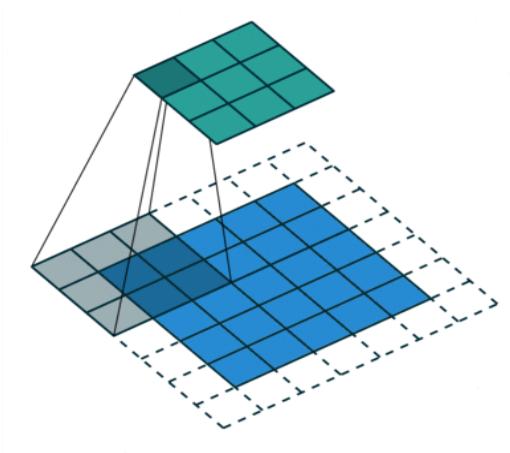
<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

## 2. Fundamentals of Deep Learning

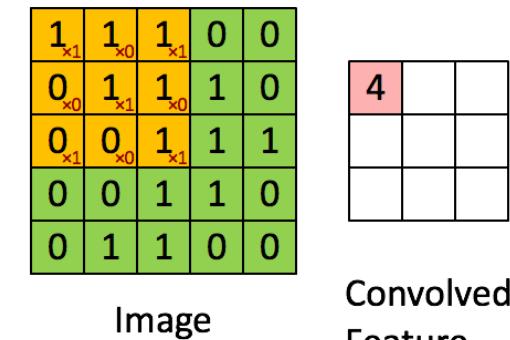
# Introduction to Convolutional Neural Networks



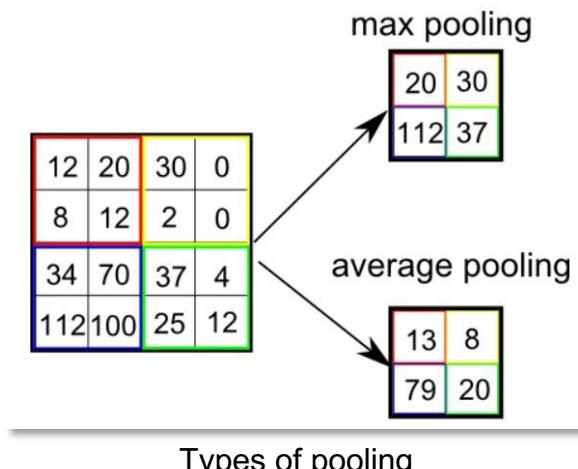
Input image



Convolution with stride=2 and padding=1



Convolving 5x5x1 image with 3x3x1 kernel to get 3x3x1 convolved feature



## 2. Fundamentals of Deep Learning

# Convolutional Neural Network Architectures – VGG (Visual Geometry Group)

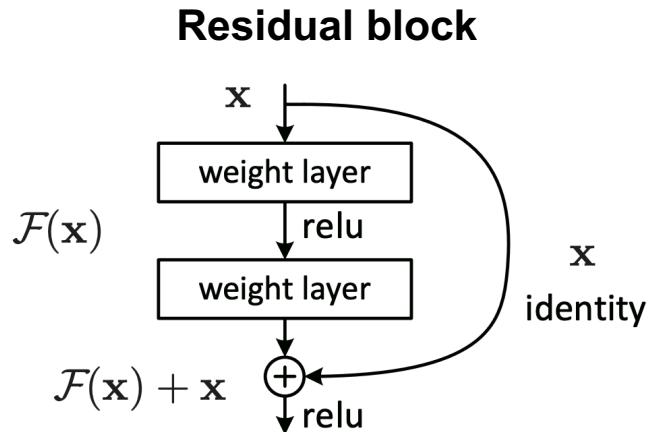
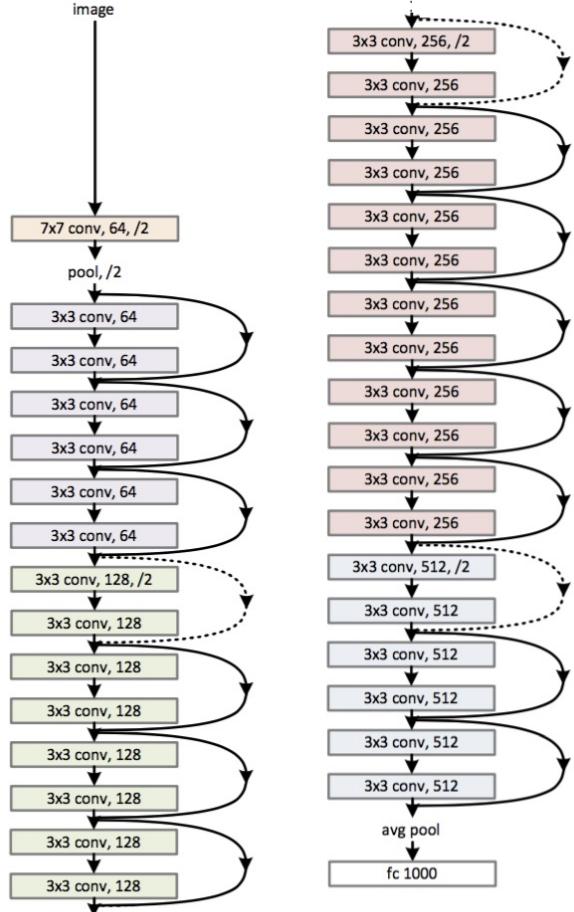
| ConvNet Configuration               |                        |                               |  |  |   |
|-------------------------------------|------------------------|-------------------------------|--|--|---|
| A                                   | A-LRN                  | B                             | C  | D  | E   |
| 11 weight layers                    | 11 weight layers       | 13 weight layers              | 16 weight layers                           | 16 weight layers                           | 19 weight layers  |
| input ( $224 \times 224$ RGB image) |                        |                               |  |  |   |
| conv3-64                            | conv3-64<br><b>LRN</b> | conv3-64<br><b>conv3-64</b>   | conv3-64<br>conv3-64                       | conv3-64<br>conv3-64                       | conv3-64<br>conv3-64                                    |
| maxpool                             |                        |                               |  |  |   |
| conv3-128                           | conv3-128              | conv3-128<br><b>conv3-128</b> | conv3-128<br>conv3-128                     | conv3-128<br>conv3-128                     | conv3-128<br>conv3-128                                  |
| maxpool                             |                        |                               |  |  |   |
| conv3-256<br>conv3-256              | conv3-256<br>conv3-256 | conv3-256<br>conv3-256        | conv3-256<br>conv3-256<br><b>conv1-256</b> | conv3-256<br>conv3-256<br><b>conv3-256</b> | conv3-256<br>conv3-256<br>conv3-256<br><b>conv3-256</b> |
| maxpool                             |                        |                               |  |  |   |
| conv3-512<br>conv3-512              | conv3-512<br>conv3-512 | conv3-512<br>conv3-512        | conv3-512<br>conv3-512<br><b>conv1-512</b> | conv3-512<br>conv3-512<br><b>conv3-512</b> | conv3-512<br>conv3-512<br>conv3-512<br><b>conv3-512</b> |
| maxpool                             |                        |                               |  |  |   |
| conv3-512<br>conv3-512              | conv3-512<br>conv3-512 | conv3-512<br>conv3-512        | conv3-512<br>conv3-512<br><b>conv1-512</b> | conv3-512<br>conv3-512<br><b>conv3-512</b> | conv3-512<br>conv3-512<br>conv3-512<br><b>conv3-512</b> |
| maxpool                             |                        |                               |  |  |   |
| FC-4096                             |                        |                               |  |  |   |
| FC-4096                             |                        |                               |  |  |   |
| FC-1000                             |                        |                               |  |  |   |
| soft-max                            |                        |                               |  |  |   |

- Main idea: increasing depth of CNNs
- $3 \times 3$  conv. layers with a stride of 1 - small receptive fields (compared to prev.  $11 \times 11$  with a stride of 4 in AlexNet)
- $1 \times 1$  conv. to make the decision function more non-linear without changing the receptive fields
- ReLU activation function
- ImageNet dataset

<https://arxiv.org/pdf/1409.1556.pdf>

## Convolutional Neural Network Architectures – ResNet

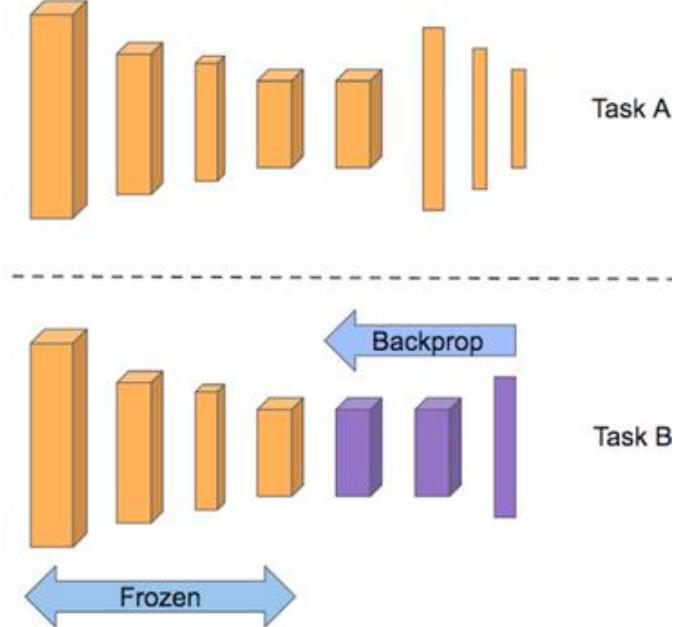
34-layer residual



- Main idea: “identity shortcut connection”
- Avoids vanishing gradient problem
- The network gets the input along with the learning on the residual and if the input function was the appropriate function, it can change the weights of the residual function to be zero.
- ImageNet dataset

<https://arxiv.org/pdf/1512.03385.pdf>

# Transfer learning – Pretrained Convolutional Neural Networks



- Pre-trained model is a saved network that was previously trained on a large dataset.
- Feature Extraction: Use the feature maps from the pre-trained model to detect features in the new samples. Add a new classifier, which will be trained from scratch to make predictions.
- Fine-Tuning: Unfreeze a few top layers of a pretrained model and jointly train both the newly-added classifier layers and the last layers of the base model. This allows us "fine-tune" the higher-order feature maps to make them more relevant for the specific task.

```
import torchvision.models as models  
model = models.resnet34(weights='IMAGENET1K_V1')
```

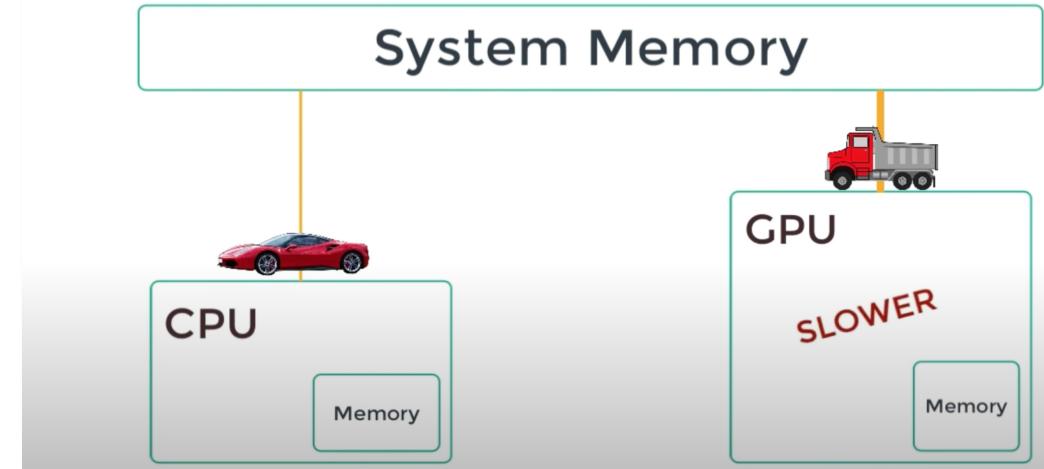
```
import torchvision.models as models  
model = models.vgg19(weights='IMAGENET1K_V1')
```

## 2. Fundamentals of Deep Learning

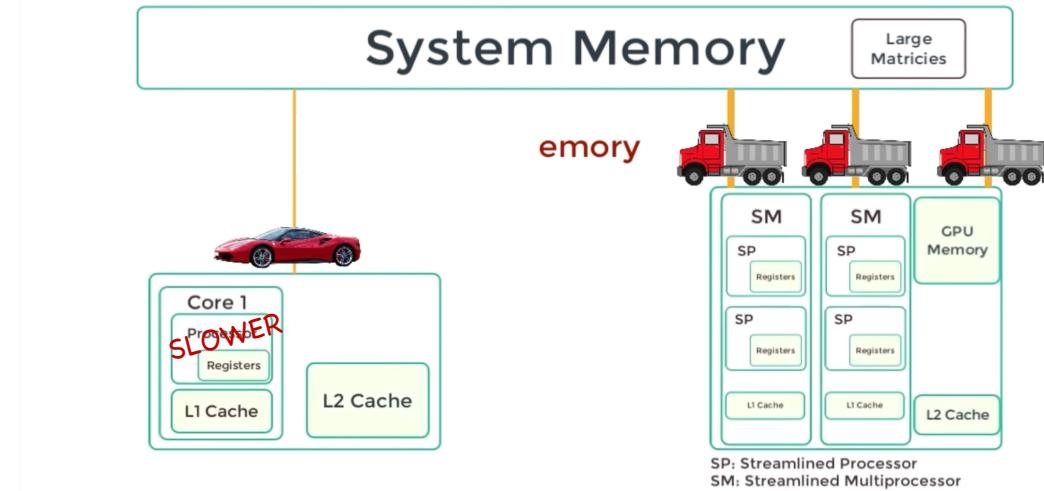
# Neural Networks and GPUs - Why?

- GPUs in comparison to CPUs:
  - GPU allows parallel running of repetitive calculations within an application
  - CPU can be thought of as the taskmaster of the entire system, coordinating a wide range of general-purpose computing tasks
  - GPU performs a narrower range of more specialized tasks (e.g., matrix multiplications)
- CPUs are faster than GPUs in scalar multiplications.
- GPUs are faster than CPUs in matrix multiplications.

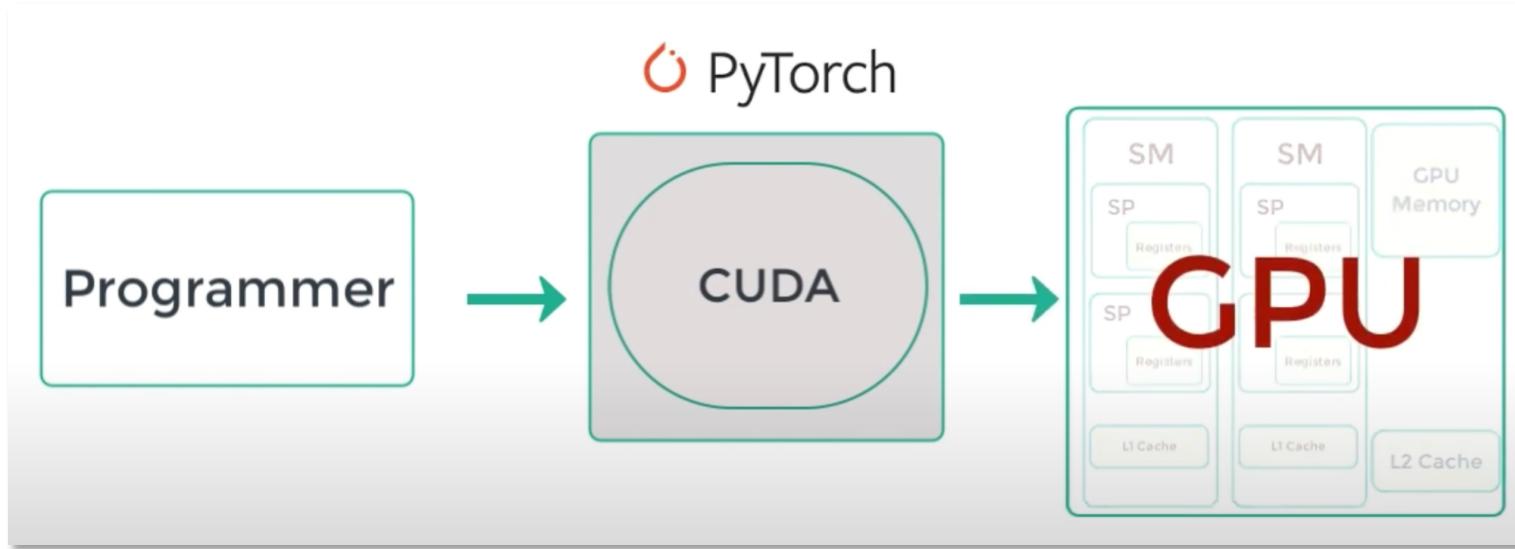
### CASE 1: Scalar Multiplication



### CASE 2: Matrix Multiplication



# Neural Networks and GPUs - How?



```
# Device configuration - cpu  
device = torch.device('cpu')  
  
# Device configuration - gpu  
device = torch.device('cuda')  
  
# Model to device  
model = model.to(device)  
  
# Data to device  
images = images.to(device=device)  
labels = labels.to(device=device)
```

# CNN Exercises in Pytorch

## 2. Fundamentals of Deep Learning

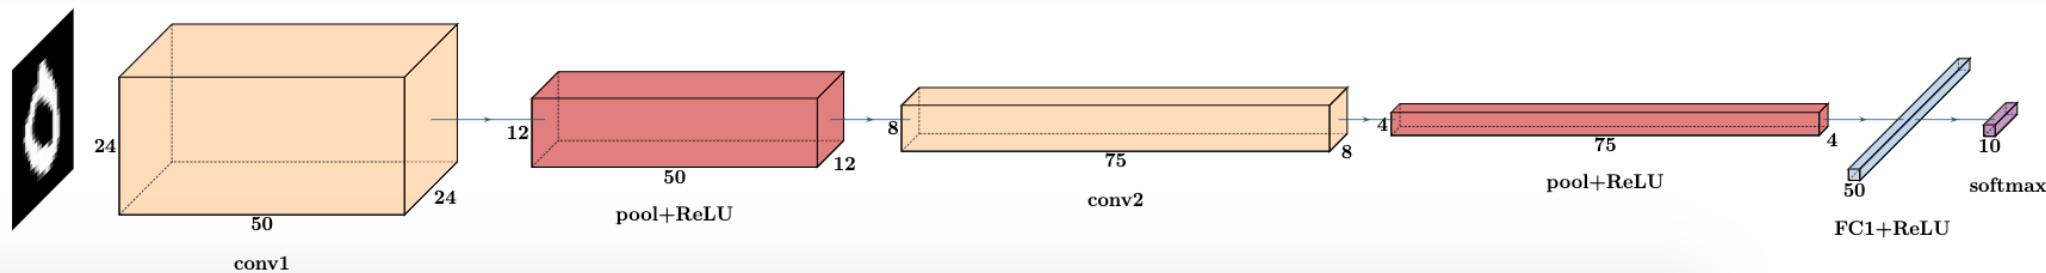
# Simple Convolutional Neural Network

```
class Model(nn.Module):
    def __init__(self, image_width, image_channels, num_classes):
        super().__init__()
        self.conv1 = nn.Conv2d(image_channels, conv1_out_channels, kernel_size)
        self.conv2 = nn.Conv2d(conv1_out_channels, conv2_out_channels, kernel_size)
        self.relu = nn.ReLU()
        self.pool = nn.MaxPool2d(pool_size)
        self.fc1 = nn.Linear(conv2_out_channels*(dim_2**2), fc1_out_channels)
        self.fc2 = nn.Linear(fc1_out_channels, num_classes)

    def forward(self, x):
        x = self.relu(self.pool(self.conv1(x)))
        x = self.relu(self.pool(self.conv2(x)))
        x = torch.flatten(x, 1)
        x = self.relu(self.fc1(x))
        x = self.fc2(x)
        return x
```

```
# Hyper-parameters
image_width = 28
image_channels = 1
conv1_out_channels = 50
conv2_out_channels = 75
kernel_size = 5
pool_size = 2
fc1_out_channels = 50
num_classes = 10
num_epochs = 5
batch_size = 100
learning_rate = 0.001
```

```
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
```

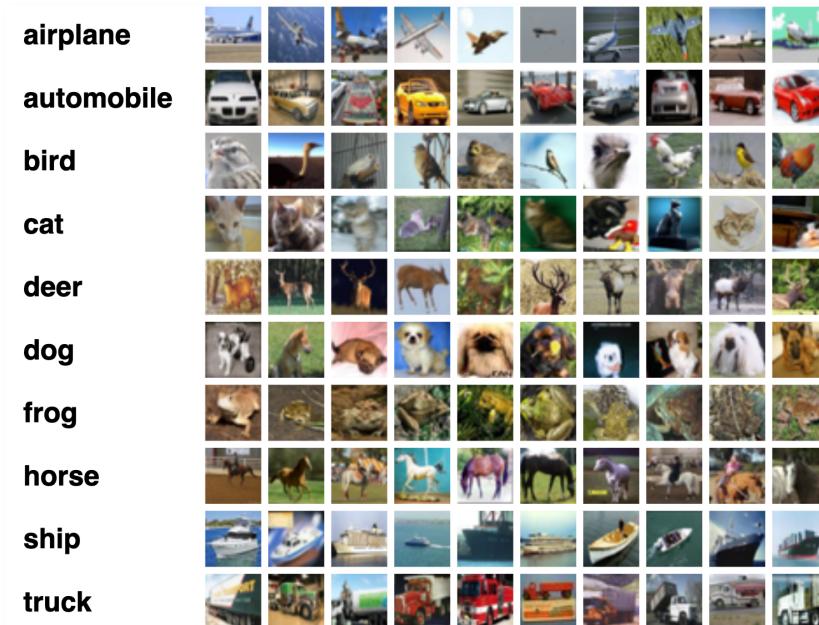


## 2. Fundamentals of Deep Learning

# CIFAR10 Data

```
# download the data
train_dataset = torchvision.datasets.CIFAR10(root='./data', train=True, transform=transforms.ToTensor(), download=True)
test_dataset = torchvision.datasets.CIFAR10(root='./data', train=False, transform=transforms.ToTensor(), download=True)

# transform to DataLoader
train_loader = torch.utils.data.DataLoader(dataset=train_dataset, batch_size=batch_size, shuffle=True)
test_loader = torch.utils.data.DataLoader(dataset=test_dataset, batch_size=batch_size, shuffle=False)
```



Input size: 32x32x3

## 2. Fundamentals of Deep Learning

### exercise1.py: Simple Convolutional Neural Network on a CPU vs on a GPU (155685 parameters)



Output:

Device: cpu

Epoch [2/20], Step [49920/50000], Loss: 1.3028  
Epoch [4/20], Step [49920/50000], Loss: 1.3503  
Epoch [6/20], Step [49920/50000], Loss: 0.9295  
Epoch [8/20], Step [49920/50000], Loss: 0.8239  
Epoch [10/20], Step [49920/50000], Loss: 0.7265  
Epoch [12/20], Step [49920/50000], Loss: 0.6835  
Epoch [14/20], Step [49920/50000], Loss: 0.8085  
Epoch [16/20], Step [49920/50000], Loss: 0.6006  
Epoch [18/20], Step [49920/50000], Loss: 0.5492  
Epoch [20/20], Step [49920/50000], Loss: 0.7049

Test set: Accuracy: 69.6 %

Elapsed time for training and testing: **396.54 s**

(396.54 s ≈ **6.6 min**)

Output:

Device: cuda

Epoch [2/20], Step [49920/50000], Loss: 1.1545  
Epoch [4/20], Step [49920/50000], Loss: 1.0863  
Epoch [6/20], Step [49920/50000], Loss: 1.0183  
Epoch [8/20], Step [49920/50000], Loss: 0.8718  
Epoch [10/20], Step [49920/50000], Loss: 0.6980  
Epoch [12/20], Step [49920/50000], Loss: 0.6620  
Epoch [14/20], Step [49920/50000], Loss: 0.6058  
Epoch [16/20], Step [49920/50000], Loss: 0.6141  
Epoch [18/20], Step [49920/50000], Loss: 0.4936  
Epoch [20/20], Step [49920/50000], Loss: 0.5196

Test set: Accuracy: 69.8 %

Elapsed time for training and testing: **130.16 s**

(130.16 s ≈ **2 min**)

## 2. Fundamentals of Deep Learning

exercise1\_resnet34.py: Pretrained ResNet34 on CIFAR10 data on a CPU vs on a GPU  
(21.8M parameters)



Output:

Device: cpu

```
Epoch [2/20], Step [49920/50000], Loss: 1.2576
Epoch [4/20], Step [49920/50000], Loss: 0.9185
Epoch [6/20], Step [49920/50000], Loss: 0.8305
Epoch [8/20], Step [49920/50000], Loss: 0.6098
Epoch [10/20], Step [49920/50000], Loss: 0.5429
Epoch [12/20], Step [49920/50000], Loss: 0.4359
Epoch [14/20], Step [49920/50000], Loss: 0.2560
Epoch [16/20], Step [49920/50000], Loss: 0.2965
Epoch [18/20], Step [49920/50000], Loss: 0.3191
Epoch [20/20], Step [49920/50000], Loss: 0.3405
```

Test set: Accuracy: 80.52 %

Elapsed time for training and testing: **2685.67 s**

(2685.67 s ≈ 45 min)

Output:

Device: cuda

```
Epoch [2/20], Step [49920/50000], Loss: 0.6397
Epoch [4/20], Step [49920/50000], Loss: 0.4903
Epoch [6/20], Step [49920/50000], Loss: 0.2941
Epoch [8/20], Step [49920/50000], Loss: 0.1125
Epoch [10/20], Step [49920/50000], Loss: 0.1250
Epoch [12/20], Step [49920/50000], Loss: 0.0748
Epoch [14/20], Step [49920/50000], Loss: 0.0519
Epoch [16/20], Step [49920/50000], Loss: 0.1285
Epoch [18/20], Step [49920/50000], Loss: 0.0511
Epoch [20/20], Step [49920/50000], Loss: 0.0567
```

Test set: Accuracy: 81.26 %

Elapsed time for training and testing: **325.72 s**

(325.72 s ≈ 5.4 min)

## 2. Fundamentals of Deep Learning

exercise1\_vgg19.py: Pretrained VGG19 on CIFAR10 data on a CPU vs on a GPU  
(143.7M parameters)



Output:

Device: cpu

```
Epoch [2/20], Step [49920/50000], Loss: 1.2591
Epoch [4/20], Step [49920/50000], Loss: 0.8687
Epoch [6/20], Step [49920/50000], Loss: 0.6782
Epoch [8/20], Step [49920/50000], Loss: 0.3366
Epoch [10/20], Step [49920/50000], Loss: 0.5239
Epoch [12/20], Step [49920/50000], Loss: 0.1858
Epoch [14/20], Step [49920/50000], Loss: 0.2304
Epoch [16/20], Step [49920/50000], Loss: 0.3175
Epoch [18/20], Step [49920/50000], Loss: 0.1343
Epoch [20/20], Step [49920/50000], Loss: 0.2626
```

Test set: Accuracy: 82.61 %

Elapsed time for training and testing: **7786.24 s**

(7786.24 s ≈ 130 min)

Output:

Device: cuda

```
Epoch [2/20], Step [49920/50000], Loss: 1.2958
Epoch [4/20], Step [49920/50000], Loss: 0.6183
Epoch [6/20], Step [49920/50000], Loss: 0.6047
Epoch [8/20], Step [49920/50000], Loss: 0.3839
Epoch [10/20], Step [49920/50000], Loss: 0.4468
Epoch [12/20], Step [49920/50000], Loss: 0.2408
Epoch [14/20], Step [49920/50000], Loss: 0.2059
Epoch [16/20], Step [49920/50000], Loss: 0.2823
Epoch [18/20], Step [49920/50000], Loss: 0.1345
Epoch [20/20], Step [49920/50000], Loss: 0.1415
```

Test set: Accuracy: 82.84 %

Elapsed time for training and testing: **383.44 s**

(383.44 s ≈ 6.4 min)

# 2. Fundamentals of Deep Learning Transformers

**Attention Is All You Need**

---

Ashish Vaswani\*  
Google Brain  
avaswani@google.com

Noam Shazeer\*  
Google Brain  
noam@google.com

Niki Parmar\*  
Google Research  
nikip@google.com

Jakob Uszkoreit\*  
Google Research  
usz@google.com

Llion Jones\*  
Google Research  
llion@google.com

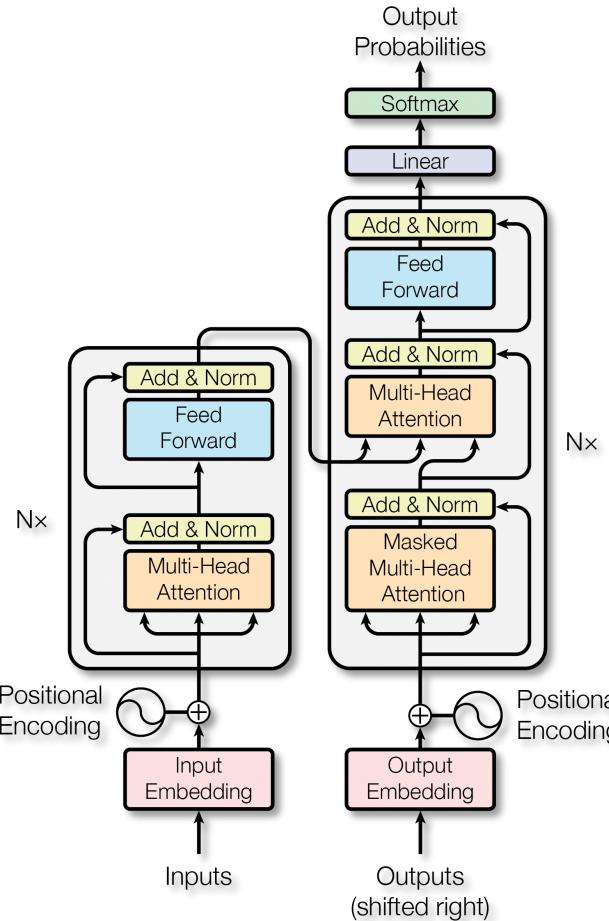
Aidan N. Gomez\*<sup>†</sup>  
University of Toronto  
aidan@cs.toronto.edu

Lukasz Kaiser\*  
Google Brain  
lukasz.kaiser@google.com

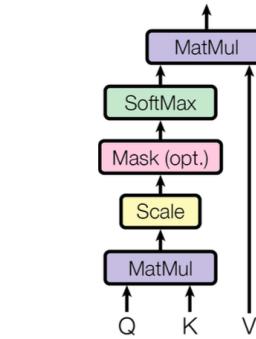
Ilia Polosukhin\*<sup>‡</sup>  
ilia.polosukhin@gmail.com

**Abstract**

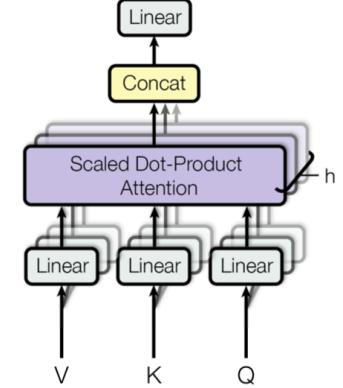
The dominant sequence transduction models are based on complex recurrent or convolutional neural networks that include an encoder and a decoder. The best performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely. Experiments on two machine translation tasks show these models to be superior in quality while being more parallelizable and requiring significantly less time to train. Our model achieves 28.4 BLEU on the WMT 2014 English-to-German translation task, improving over the existing best results, including ensembles, by over 2 BLEU. On the WMT 2014 English-to-French translation task, our model establishes a new single-model state-of-the-art BLEU score of 41.8 after training for 3.5 days on eight GPUs, a small fraction of the training costs of the best models from the literature. We show that the Transformer generalizes well to other tasks by applying it successfully to English constituency parsing both with large and limited training data.



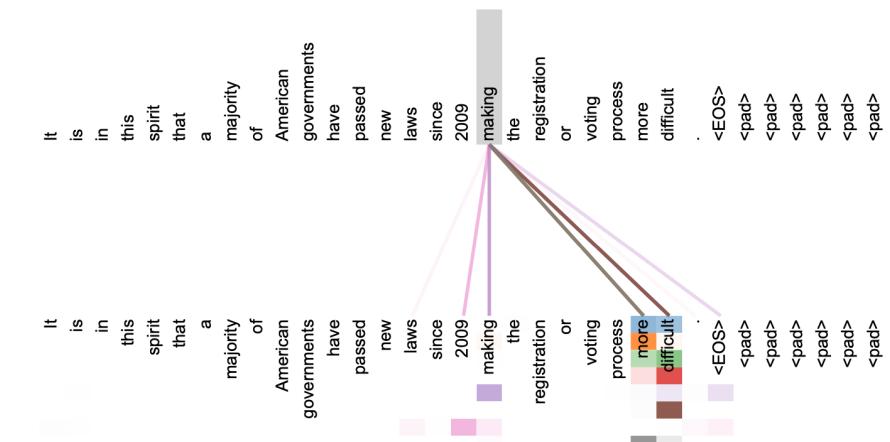
Scaled Dot-Product Attention



Multi-Head Attention



Attention Visualizations



## 2. Fundamentals of Deep Learning Transformer Exercise

```
class TransformerModel(nn.Module):
    def __init__(self, ntoken, d_model, nhead, d_hid, nlayers, dropout):
        super().__init__()
        self.encoder = nn.Embedding(ntoken, d_model)
        self.d_model = d_model
        self.pos_encoder = PositionalEncoding(d_model, dropout)
        encoder_layers = nn.TransformerEncoderLayer(d_model, nhead, d_hid, dropout)
        self.transformer_encoder = nn.TransformerEncoder(encoder_layers, nlayers)
        self.decoder = nn.Linear(d_model, ntoken)

    def forward(self, src, src_mask):
        """ src: Tensor, shape ``[seq_len, batch_size]``
            src_mask: Tensor, shape ``[seq_len, seq_len]``
            output: Tensor, shape ``[seq_len, batch_size, ntoken]`` """
        src = self.encoder(src) * math.sqrt(self.d_model)
        src = self.pos_encoder(src)
        output = self.transformer_encoder(src, src_mask)
        output = self.decoder(output)
        return output
```

- Training a sequence-to-sequence model based on **nn.Transformer**.
- The nn.Transformer module is modularized – a single component, like TransformerEncoder can be easily adapted / composed.
- We train a **nn.TransformerEncoder** to assign a probability for the likelihood of a given word (or a sequence of words) to follow a sequence of words.
- Square attention mask is required with the input sequence, because the self-attention layers are only allowed to see the earlier positions in the sequence.
- PositionalEncoding class is defined to use sine and cosine functions of different frequencies to inject the information on the position of the tokens in the input sequence.

## 2. Fundamentals of Deep Learning

# WikiText-2 Data

- Using **torchtext** to generate **WikiText-2** dataset.
- WikiText language modeling dataset is a collection of tokens extracted from the set of verified good and featured Wikipedia articles.
- Number of lines per split:
  - train 36718, valid 3760, test 4358.
  - Given a 1-D vector of sequential data, function **batchify()** arranges it into batch\_size columns and trims to fit. Batching enables more parallelizable processing, however the model treats each column independently, e.g., the dependence of G and F on the example on the right.
  - Function **get\_batch()** generates a pair of input-target sequences for the transformer model. It subdivides the source data into chunks of length bptt, e.g., with bptt=2, we'd get the input and target as shown on the right.

Dataset Preview

Subset: wikitext-2-v1 (44.8k rows)

Split: train (36.7k rows)

text (string)

```
..  
" = Valkyria Chronicles III = "  
..  
" Senjō no Valkyria 3 : cunk> Chronicles ( Japanese : 戦場のヴァルキリア3 , lit . Valkyria of the Battlefield 3 ) , commonly referred to as Valkyria Chronicles III outside Japan , is a tactical role @-@ playing video game developed by Sega and Media Vision for the PlayStation Portable . Released in January 2011 in Japan , it is the third game in the Valkyria series . cunk> the same fusion of tactical and real @-@ time gameplay as its predecessors , the story runs parallel to the first game and follows the " Nameless " , a penal military unit serving the nation of Gallia during the Second European War who perform secret black operations and are pitted against the Imperial unit " cunk> Raven " . "  
" The game began development in 2010 , carrying over a large portion of the work done on Valkyria Chronicles II . While it retained the standard features of the series , it also underwent multiple adjustments , such as making the game more cunk> for series newcomers . Character designer cunk> Honjou and composer Hitoshi Sakimoto both returned from previous entries , along with Valkyria Chronicles II ."  
" It met with positive sales in Japan , and was praised by both Japanese and western critics . After release , it received downloadable content , along with an expanded edition in November of that year . It was also adapted into manga and an original video animation series . Due to low sales of Valkyria Chronicles II , Valkyria Chronicles III was not localized , but a fan translation compatible with the game . "  
..  
" == Gameplay == "
```

$$\begin{bmatrix} A & B & C & \dots & X & Y & Z \end{bmatrix} \Rightarrow \begin{bmatrix} A \\ B \\ C \\ D \\ E \\ F \end{bmatrix} \quad \begin{bmatrix} G \\ H \\ I \\ J \\ K \\ L \end{bmatrix} \quad \begin{bmatrix} M \\ N \\ O \\ P \\ Q \\ R \end{bmatrix} \quad \begin{bmatrix} S \\ T \\ U \\ V \\ W \\ X \end{bmatrix}$$

| <b>Input</b>                                  | <b>Target</b>                                 |
|---|---|
| $\begin{bmatrix} A & G & M & S \end{bmatrix}$ | $\begin{bmatrix} B & H & N & T \end{bmatrix}$ |
| $\begin{bmatrix} B & H & N & T \end{bmatrix}$ | $\begin{bmatrix} C & I & O & U \end{bmatrix}$ |

## 2. Fundamentals of Deep Learning

### exercise1\_nlp.py: Transformer on a CPU vs on a GPU (77,879,406)



- Embedding dimension of 1024, hidden size of 1024, 4 attention heads and 3 transformer layers.

Output:

Device: cpu

```
| end of epoch 1 | time: 1618.58s | valid loss 5.75 | valid ppl 313.30  
| end of epoch 2 | time: 1635.81s | valid loss 5.69 | valid ppl 296.32  
| end of epoch 3 | time: 1493.84s | valid loss 5.57 | valid ppl 263.73  
| End of training | test loss 5.48 | test ppl 240.33
```

(4748.23 s ≈ 80 min)

Output:

Device: cuda

```
| end of epoch 1 | time: 79.06s | valid loss 5.75 | valid ppl 314.83  
| end of epoch 2 | time: 70.25s | valid loss 5.58 | valid ppl 265.83  
| end of epoch 3 | time: 69.92s | valid loss 5.49 | valid ppl 242.84  
| End of training | test loss 5.40 | test ppl 221.61
```

(219.23 s ≈ 3.7 min)

## Exercise: Start an interactive job and create an extended Enroot image



1. Extend the existing Pytorch container with the following python packages:
  - pip install --no-cache-dir torchdata portalocker torchvision  
(for exercise1\_nlp.py, Transformer exercises)
  - HOROVOD\_GPU\_OPERATIONS=NCCL pip install --no-cache-dir horovod  
(for exercise2.py, Horovod exercise which will be shown later in the slides)
2. Run exercises `exercise1_resnet34.sbatch`, `exercise1_vgg19.sbatch`, and/or `exercise1_nlp.sbatch` as batch jobs.
3. Run exercise `exercise1_nlp.ipynb` as a Jupyter Notebook on <https://datalab3.srv.lrz.de/>.

## 1. Introduction to the LRZ AI Systems

.....

- ✓ Overview of the LRZ AI Systems
- ✓ Access to the LRZ AI Systems
- ✓ NVIDIA NGC Cloud
- ✓ Introduction to Enroot Containers
- ✓ Interactive and Batch Jobs
- ✓ Open on Demand
- ✓ Exercise: Run a job with an Enroot container

## 2. Fundamentals of Deep Learning

.....

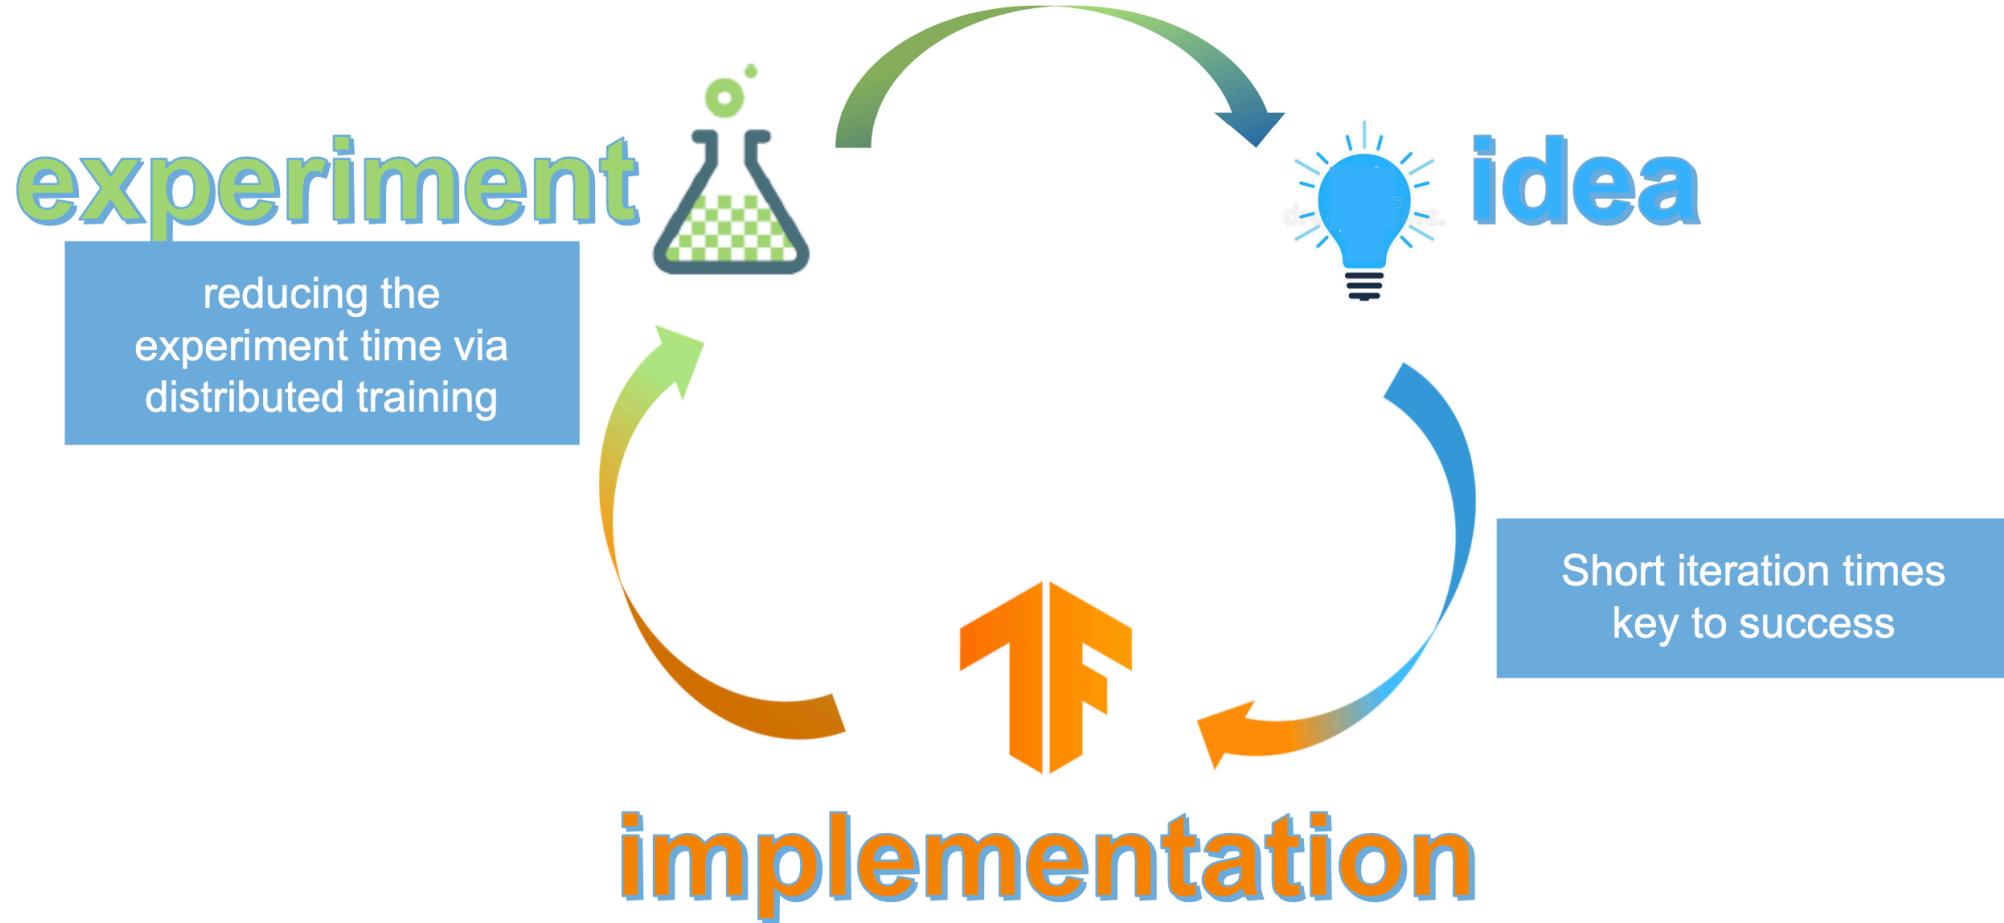
- ✓ Introduction to Neural Networks
- ✓ Introduction to Convolutional Neural Networks
- ✓ Exercises: Train CNNs on a GPU
- ✓ Introduction to Transformers
- ✓ Exercise: Train a Transformer on a GPU

## 3. Distributed Training of Neural Networks

.....

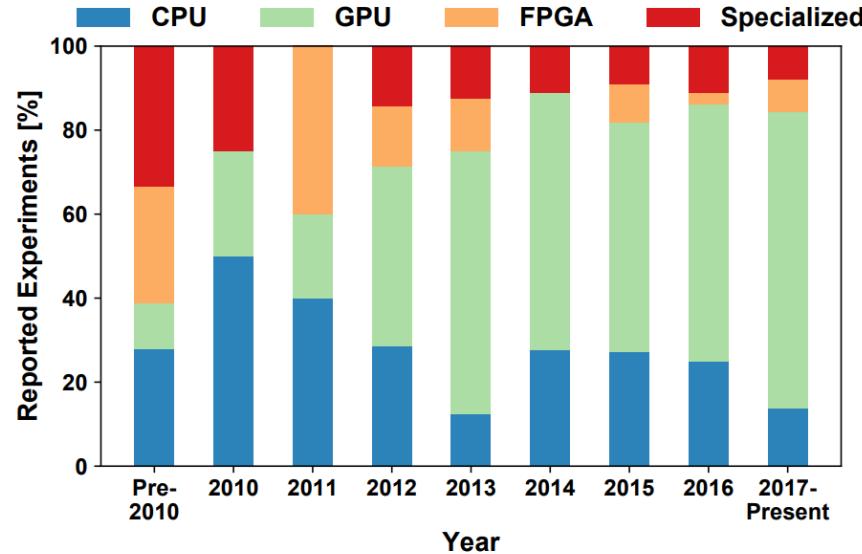
- Data Parallel Training
- Exercise: DP Training of CNN on 2 GPUs
- Model Parallel Training: Pipeline Parallel and Tensor Parallel
- Exercise: PP Training of Transformer on 2 GPUs

## Experimental Science Requires Short Iteration Times

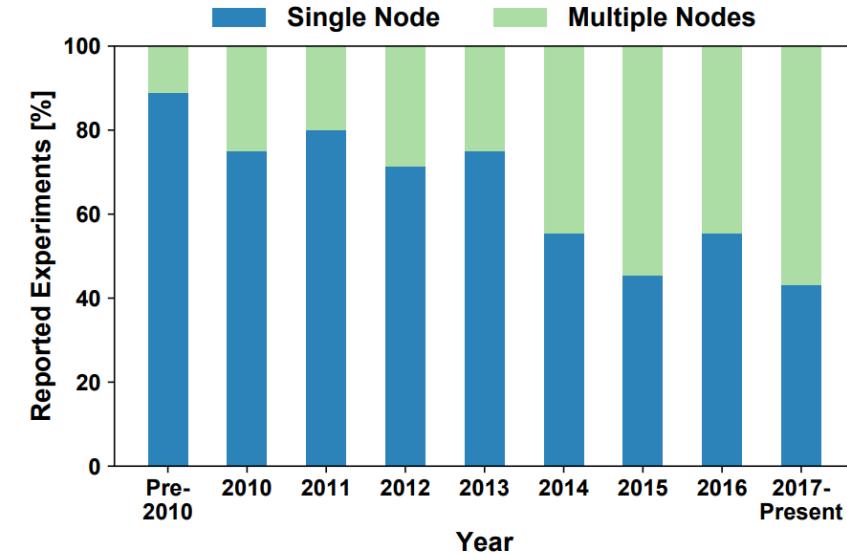


### 3. Distributed Training of Neural Networks

## Trends of Parallel Architectures in Deep Learning



(a) Hardware Architectures



(b) Training with Single vs. Multiple Nodes

Fig. 3. Parallel Architectures in Deep Learning

Source: Ben-Nun and Hoefer, 2018, arXiv:1802.09941

### 3. Distributed Training of Neural Networks

## Increasing Amount of Data Available for Deep Learning

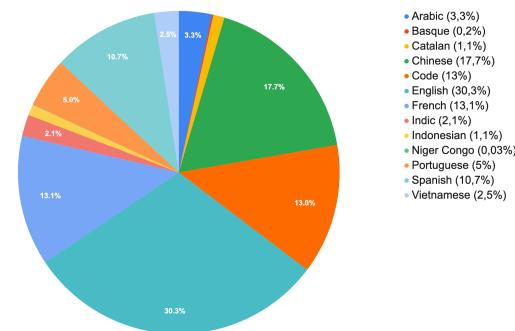
ImageNet



14,197,122 images  
(150 GB)

<https://paperswithcode.com/dataset/imagenet>

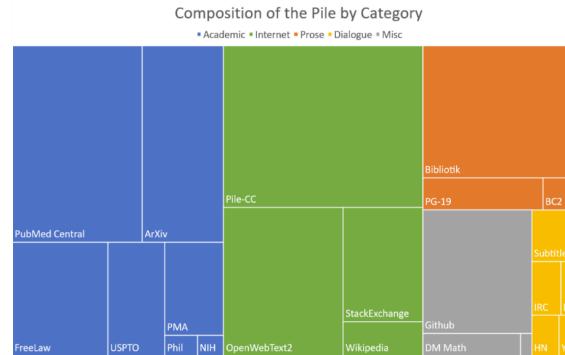
BigScience  
Multilingual Dataset  
for Language  
Modeling



350 billion tokens  
(1.5 TB)  
46 languages

<https://bigscience.huggingface.co/blog/building-a-tb-scale-multilingual-dataset-for-language-modeling>

Pile: Dataset of  
Diverse Text for  
Language Modeling



800GB

<https://arxiv.org/abs/2101.00027>

Common Crawl

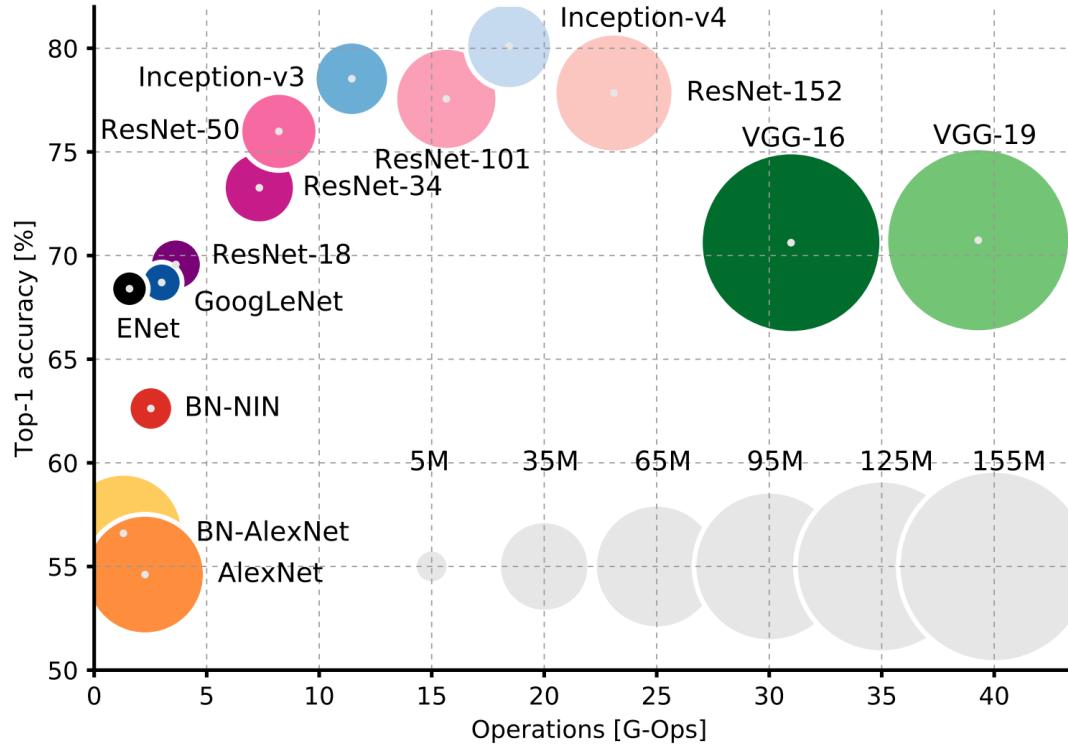


2.95 billion web pages  
(351.844 TB)

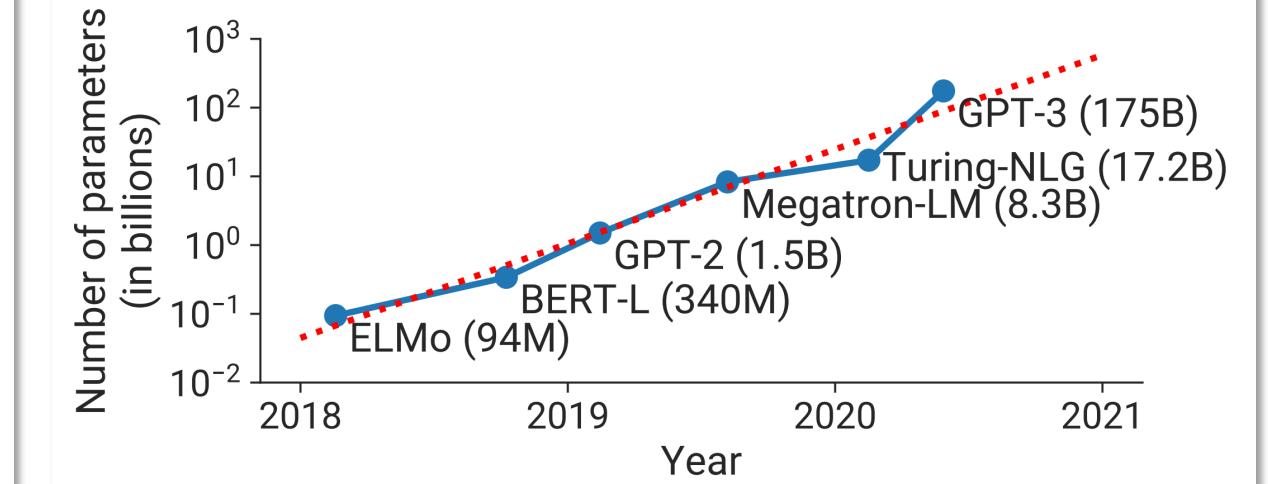
<https://commoncrawl.org/>

### 3. Distributed Training of Neural Networks

## Increasing Complexity of Deep Learning Models



Source: Canziani et al., 2017, arXiv:1605.07678



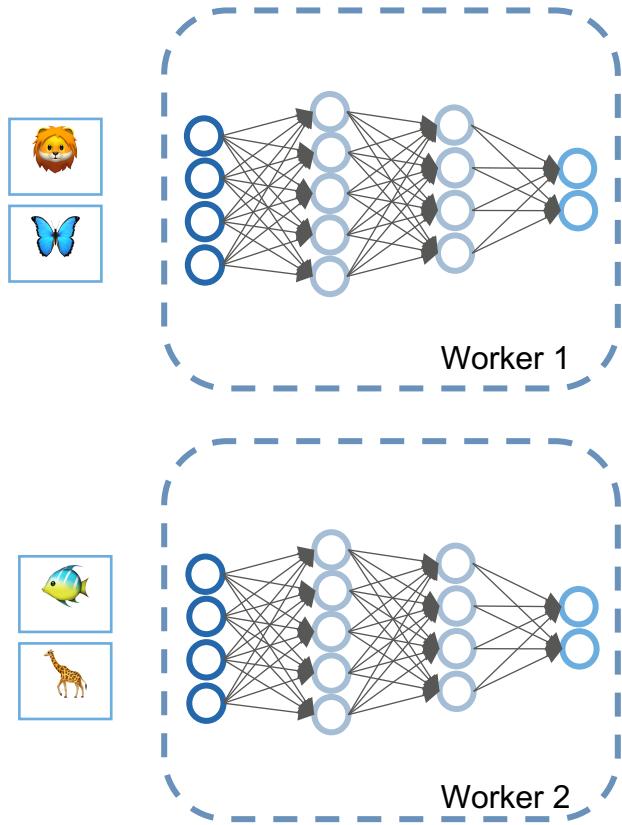
Source: Narayanan et al., 2021, arXiv:2104.04473

BigScience's BLOOM – 176B  
Google's PaLM – 540B  
GPT-4 – reportedly 1T

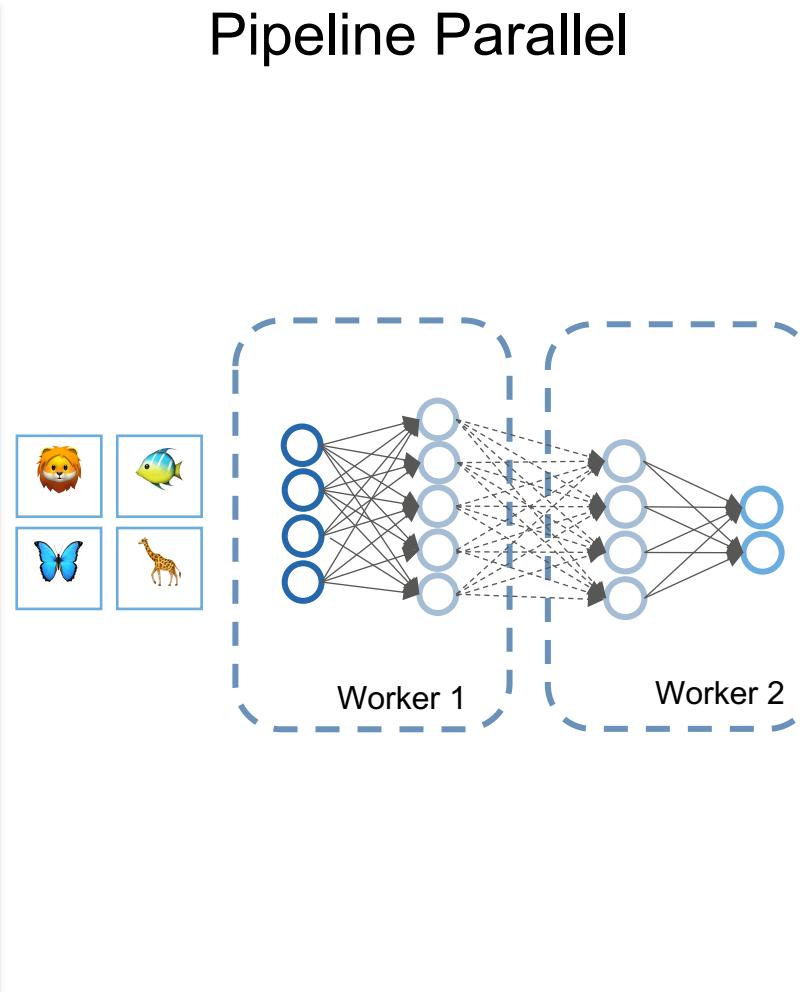
### 3. Distributed Training of Neural Networks

## Overview of the Techniques for Distributed Training of NNs

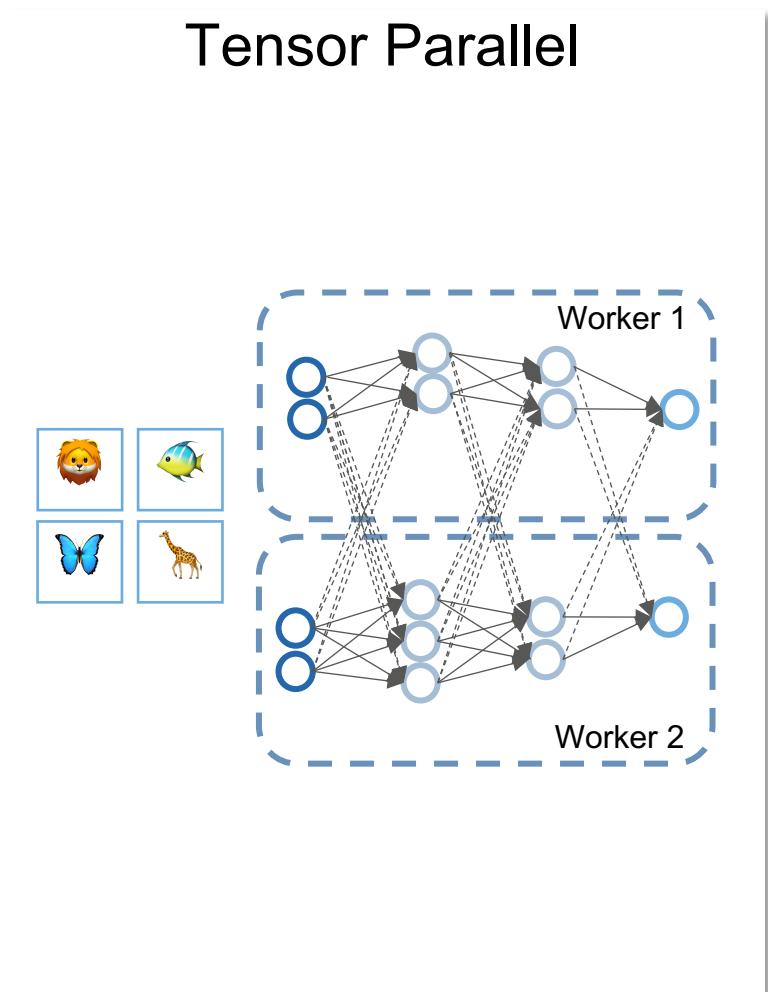
Data Parallel



Pipeline Parallel



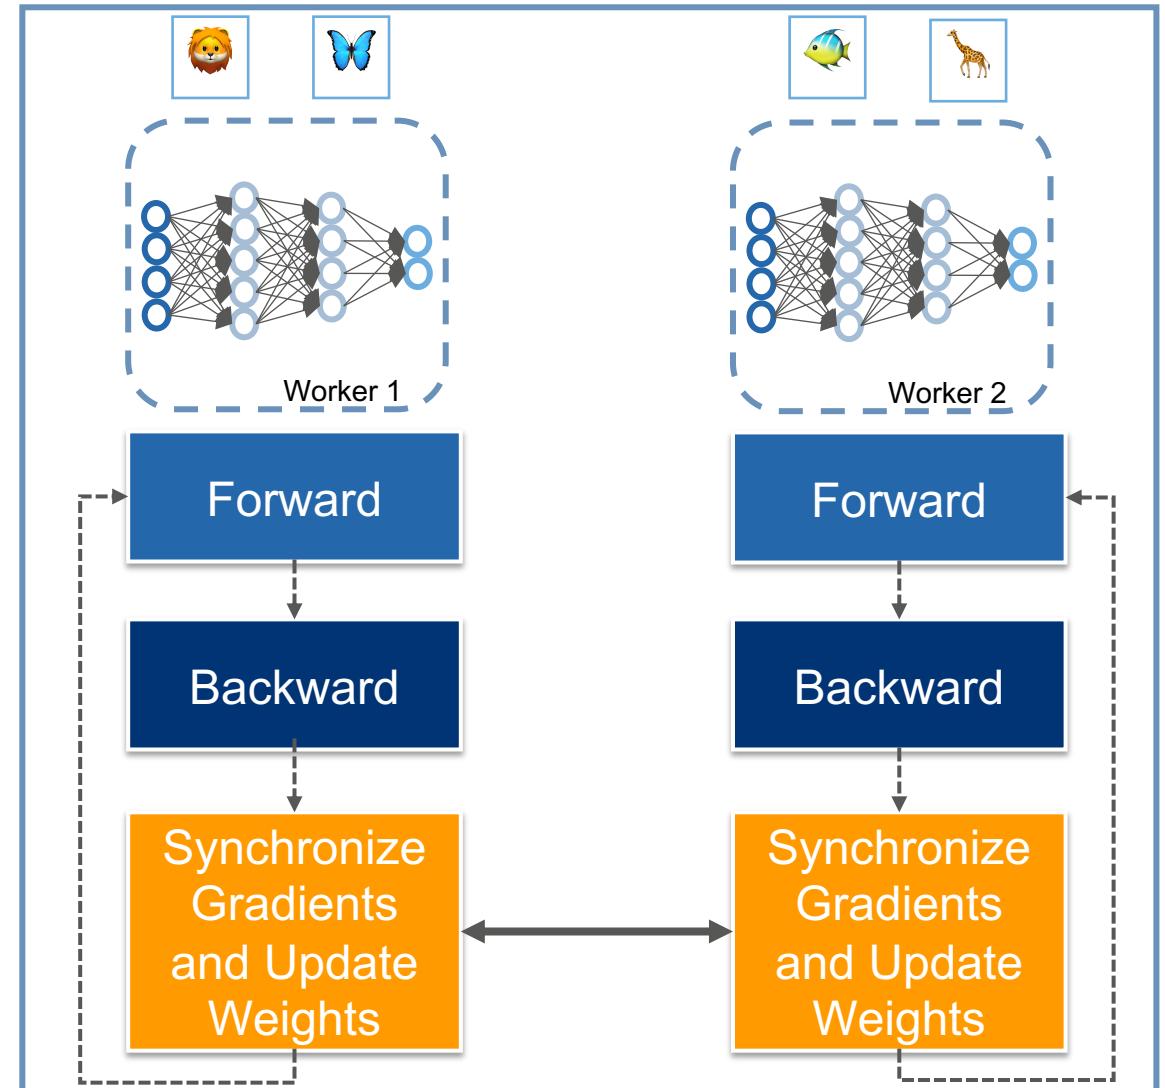
Tensor Parallel



### 3. Distributed Training of Neural Networks

## Data Parallel Training of Neural Networks

- Every worker has a copy of the whole model and a subset of the data.
- Each worker generates its own weights, biases, gradients and optimizer states.
- At the end of every iteration, workers aggregate the gradients across all workers, average gradients, and then apply those averaged gradients (using e.g. *allreduce*).
- **Collective (all-to-all)** communication of weight gradients.



### 3. Distributed Training of Neural Networks

#### Horovod

- Distributed open-source deep learning training framework for TensorFlow, Keras, PyTorch, and Apache MXNet.
- Runs across *multiple GPUs* and across *multiple machines/nodes*.
- *Fast* – scales up to hundreds of GPUs with upwards of 90% scaling efficiency.
- *Easy* – a few lines of codes.
- *Portable* – different frameworks.



#1 initialization  
import horovod.torch as hvd  
hvd.init()

#2 pin each GPU to a single process  
if torch.cuda.is\_available():  
 torch.cuda.set\_device(hvd.local\_rank())

#3 distributed optimizer  
opt = #choose your optimizer of preference  
opt = hvd.DistributedOptimizer(opt)

#4 broadcast the initial variable states from rank 0 to all other processes  
hvd.broadcast\_parameters(model.state\_dict(), root\_rank=0)  
hvd.broadcast\_optimizer\_state(optimizer, root\_rank=0)

#5 differentiate among different workers (e.g. check pointing on worker 0)  
if hvd.rank() == 0:  
 checkpoint.save(checkpoint\_dir)

```
# run training with 2 GPUs on a single machine
$ horovodrun -np 2 python train.py
# run training with 2 GPUs on two machines (1 GPUs each)
$ horovodrun -np 2 -H hostname1:1,hostname2:1 python train.py
```

### 3. Distributed Training of Neural Networks

## exercise2.py: Data Parallel Training of ResNet34 on 2 GPUs with Horovod



```
# Initialize Horovod and pin each GPU to a single process
hvd.init()
if torch.cuda.is_available():
    torch.cuda.set_device(hvd.local_rank())

# Partition datasets among workers using DistributedSampler
sampler = DistributedSampler(data, num_replicas=hvd.size(), rank=hvd.rank())
loader = torch.utils.data.DataLoader(data, batch_size, sampler)

# Use distributed optimizer and broadcast parameters
optimizer = hvd.DistributedOptimizer(optimizer, model.named_parameters())
hvd.broadcast_parameters(model.state_dict(), root_rank=0)
hvd.broadcast_optimizer_state(optimizer, root_rank=0)

# Horovod: average metric values across workers.
def metric_average(val, name):
    tensor = val.clone().detach()
    avg_tensor = hvd.allreduce(tensor, name=name)
    return avg_tensor.item()
```

- Install in Pytorch container with:  
HOROVOD\_GPU\_OPERATIONS=NCCL pip install --no-cache-dir horovod

```
$ horovodrun -np 2 python exercise2.py
```

Output:

```
[1,1]<stdout>:Epoch [2/10], Step [24960/25000], Loss: 0.5788
[1,0]<stdout>:Epoch [2/10], Step [24960/25000], Loss: 0.4635
[1,1]<stdout>:Epoch [4/10], Step [24960/25000], Loss: 0.3478
[1,0]<stdout>:Epoch [4/10], Step [24960/25000], Loss: 0.2552
[1,1]<stdout>:Epoch [6/10], Step [24960/25000], Loss: 0.2159
[1,0]<stdout>:Epoch [6/10], Step [24960/25000], Loss: 0.1682
[1,1]<stdout>:Epoch [8/10], Step [24960/25000], Loss: 0.1184
[1,0]<stdout>:Epoch [8/10], Step [24960/25000], Loss: 0.0652
[1,0]<stdout>:Epoch [10/10], Step [24960/25000], Loss: 0.1116
[1,1]<stdout>:Epoch [10/10], Step [24960/25000], Loss: 0.0584

[1,0]<stdout>:Test set: Accuracy: 81.08%
[1,0]<stdout>:Elapsed time: 126.31 s ≈ 2 min
```

### 3. Distributed Training of Neural Networks

## Pytorch Distributed Data Parallel

- Based on package *torch.distributed* for synchronizing gradients.
- DDP registers a hook for each parameter that fires when the corresponding gradient is computed in the backward pass. That signal is used to trigger gradient synchronization across processes.
- Runs across *multiple GPUs* and across *multiple machines/nodes*.
- Near-linear scalability using 256 GPUs.

Li et al., 2020, arXiv:2006.15704

<https://pytorch.org/docs/master/notes/ddp.html>

[https://pytorch.org/tutorials/intermediate/ddp\\_tutorial.html#basic-use-case](https://pytorch.org/tutorials/intermediate/ddp_tutorial.html#basic-use-case)

```
import torch.distributed as dist
import torch.multiprocessing as mp
from torch.nn.parallel import DistributedDataParallel as DDP

def example(rank, world_size):
    # create default process group – nccl for building with cuda
    dist.init_process_group(backend="nccl", rank=rank,
                           world_size=world_size)

    # create a model and wrap it with DDP
    model = nn.Linear(10, 10).to(rank)
    ddp_model = DDP(model, device_ids=[rank])
    # define loss function and optimizer
    loss_fn = nn.MSELoss()
    optimizer = torch.optim.SGD(ddp_model.parameters(), lr=0.01)

    # forward pass
    outputs = ddp_model(torch.randn(20, 10).to(rank))
    labels = torch.randn(20, 10).to(rank)
    # backward pass and update parameters
    loss_fn(outputs, labels).backward()
    optimizer.step()

    # cleanup
    dist.destroy_process_group()

if __name__ == "__main__":
    # environment variables for using torch.distributed
    os.environ["MASTER_ADDR"] = "localhost"
    os.environ["MASTER_PORT"] = "29500"
    world_size = 2 # world_size = gpus * nodes
    mp.spawn(example, args=(world_size,), nprocs=world_size)
```

### 3. Distributed Training of Neural Networks

## exercise3.py: Data Parallel Training of VGG19 on 2 GPUs with Pytorch DDP



```
# environment variables for using torch.distributed
world_size = gpus * nodes
os.environ['MASTER_ADDR'] = 'localhost'
os.environ['MASTER_PORT'] = '8888'
mp.spawn(train_and_test, gpus, args=(gpus, nodes, nr, world_size))

def train_and_test(gpu, gpus, nodes, nr, world_size):
    rank = nr * gpus + gpu # nr is ranking within the nodes
    dist.init_process_group(backend='nccl', init_method='env://', world_size, rank)
    # Set the seed to initiate the model with the same weights on each worker
    torch.manual_seed(0)
    # Set the model and loss to GPU
    model.cuda()
    criterion = nn.CrossEntropyLoss().cuda()
    # Wrap the model
    model = nn.parallel.DistributedDataParallel(model, device_ids=[gpu])
    # Partition datasets among workers using DistributedSampler
    sampler = DistributedSampler(data, num_replicas=world_size, rank=rank)
    loader = torch.utils.data.DataLoader(data, batch_size, sampler)
```

\$ python exercise3.py

Output:

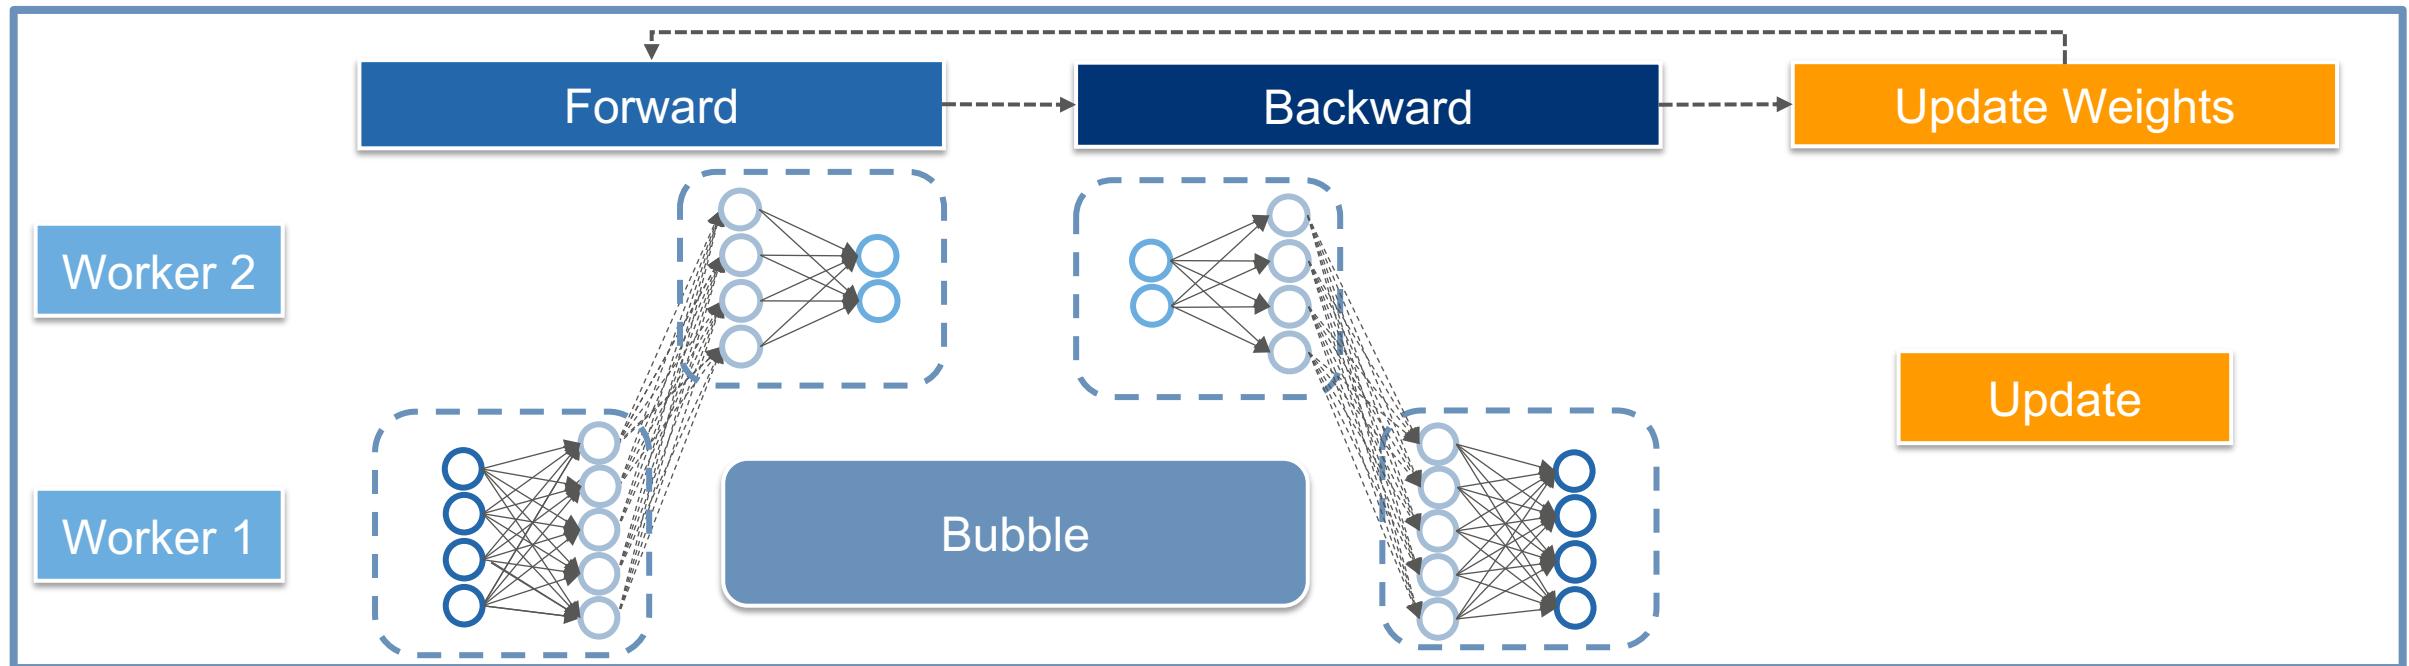
```
Epoch [2/10], Step [24960/25000], Loss: 1.5240
Epoch [2/10], Step [24960/25000], Loss: 1.4402
Epoch [4/10], Step [24960/25000], Loss: 0.7013
Epoch [4/10], Step [24960/25000], Loss: 0.7268
Epoch [6/10], Step [24960/25000], Loss: 0.5393
Epoch [6/10], Step [24960/25000], Loss: 0.4175
Epoch [8/10], Step [24960/25000], Loss: 0.3391
Epoch [8/10], Step [24960/25000], Loss: 0.4450
Epoch [10/10], Step [24960/25000], Loss: 0.3889
Epoch [10/10], Step [24960/25000], Loss: 0.3106
```

Test set: Accuracy: 81.76%

Elapsed time: **234.65 s** ≈ 4 min

## Pipeline Parallel Training of Neural Networks

- Each worker has one or more layers, and the corresponding parameters of the neural network. There is a single copy of weight parameters shared between the workers.
- Goal: Make feasible models, that are larger than the memory capacity of a single GPU.
- **Point-to-point communication of intermediate activations and gradients.**



### 3. Distributed Training of Neural Networks

## Pipeline Parallel Training of Neural Networks - Different Schemes

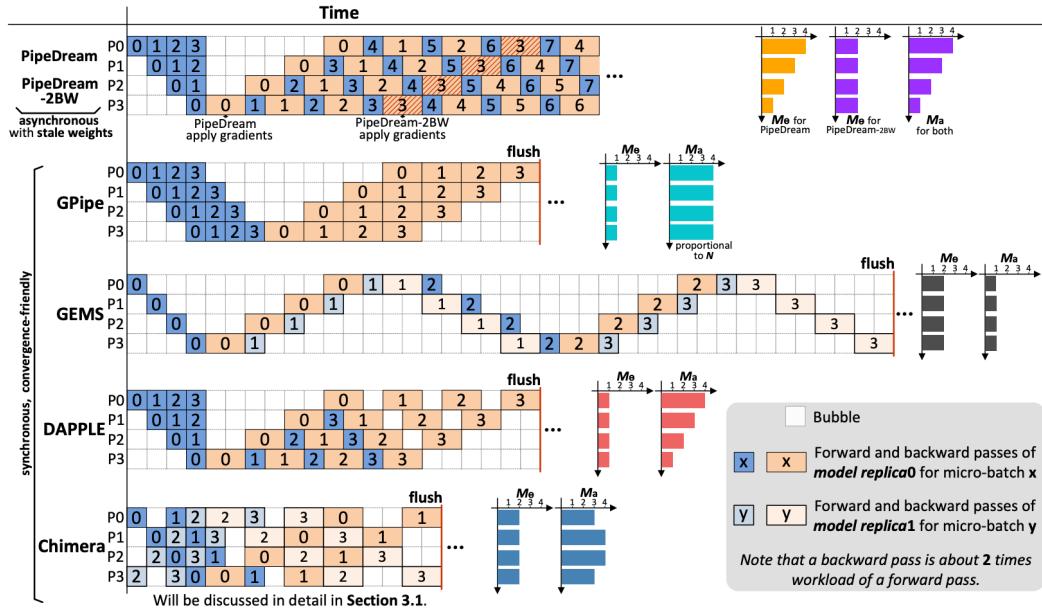


Figure 2: Pipeline parallelism schemes, with four pipeline stages ( $D=4$ ) and four micro-batches ( $N=4$ ) within a training iteration, except that PipeDream updates the model after each backward pass on a micro-batch.

- Goal: Minimizing bubbles in the pipeline.

Li and Hoefer, 2021,  
<http://htor.inf.ethz.ch/publications/img/shigang-chimera.pdf>

Table 1: The list of symbols frequently used in the paper.

|            |   |
|------------|---|
| $D$        | The number of pipeline stages ( <i>depth</i> )  |
| $W$        | The number of replicated pipelines ( <i>width</i> ) for data parallelism <sup>1</sup> |
| $P$        | The number of workers ( $= W * D$ )   |
| $B$        | Micro-batch size  |
| $N$        | The number of micro-batches executed by each worker within a training iteration       |
| $\hat{B}$  | Mini-batch size ( $= B * N * W$ )   |
| $M_\theta$ | Memory consumption for the weights of one stage                                       |
| $M_a$      | Memory consumption for the activations of one stage                                   |

<sup>1</sup> This paper considers the cases where all pipeline stages have balanced workload, and therefore are equally replicated to combine with data parallelism.

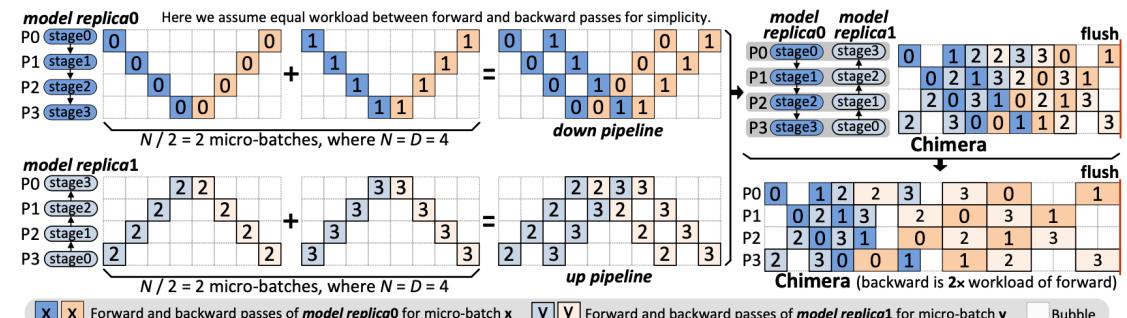


Figure 3: Model replicas and bidirectional pipelines scheduling of Chimera.

## Pipeline Parallelism with Pytorch Pipe

- Micro-batch pipeline parallelism based on Gpipe scheme.
- Depends on the Remote Procedure Call (RPC) framework from *torch.distributed* for communication.
- Currently runs only on a *single machine/node*.
- *Experimental and subject to change! Not yet fully compatible with CUDA!*

<https://pytorch.org/docs/stable/pipeline.html>

Kim et al., 2020, <https://arxiv.org/abs/2004.09910>

```
from torch.distributed import rpc
from torch.distributed.pipeline.sync import Pipe

# initialize Remote Procedure Call (RPC) framework
os.environ['MASTER_ADDR'] = "localhost"
os.environ['MASTER_PORT'] = "29500"
rpc.init_rpc(name='worker', rank=0, world_size=1)

# build the pipe with each layer on different GPU
fc1 = nn.Linear(16, 8).cuda(0)
fc2 = nn.Linear(8, 4).cuda(1)
model = nn.Sequential(fc1, fc2)
model = Pipe(model, chunks=8)
input = torch.rand(16, 16).cuda(0)

# rref is a reference to a value of a tensor on a remote worker
output_rref = model(input)
outputs = output_rref.local_value()
loss = criterion(outputs, targets.cuda(1))
```

### 3. Distributed Training of Neural Networks

## Exercise: Pipeline Parallel Training of a Transformer on 2 GPUs

```
nlayers = 12
num_gpus = 2
partition_len = ((nlayers - 1) // num_gpus) + 1
# Add encoder in the beginning.
tmp_list = [Encoder(ntokens, emsize, dropout).cuda(0)]
module_list = []
# Add all the necessary transformer blocks.
for i in range(nlayers):
    transformer_block = TransformerEncoderLayer(emsize, nhead, nhid, dropout)
    if i != 0 and i % (partition_len) == 0:
        module_list.append(nn.Sequential(*tmp_list))
        tmp_list = []
    device = i // (partition_len)
    tmp_list.append(transformer_block.to(device))
# Add decoder in the end.
tmp_list.append(Decoder(ntokens, emsize).cuda(num_gpus - 1))
module_list.append(nn.Sequential(*tmp_list))

from torch.distributed.pipeline.sync import Pipe
# Build the pipeline.
chunks = 8
model = Pipe(torch.nn.Sequential(*module_list), chunks = chunks)
```

- Using an embedding dimension of 4096, hidden size of 4096, 16 attention heads and 12 transformer layers.
- Model with ~1.4 billion parameters.
- Split the model such that half of the nn.TransformerEncoderLayer is on 1<sup>st</sup> GPU and the other half is on the 2<sup>nd</sup>.

\$ python exercise3.py

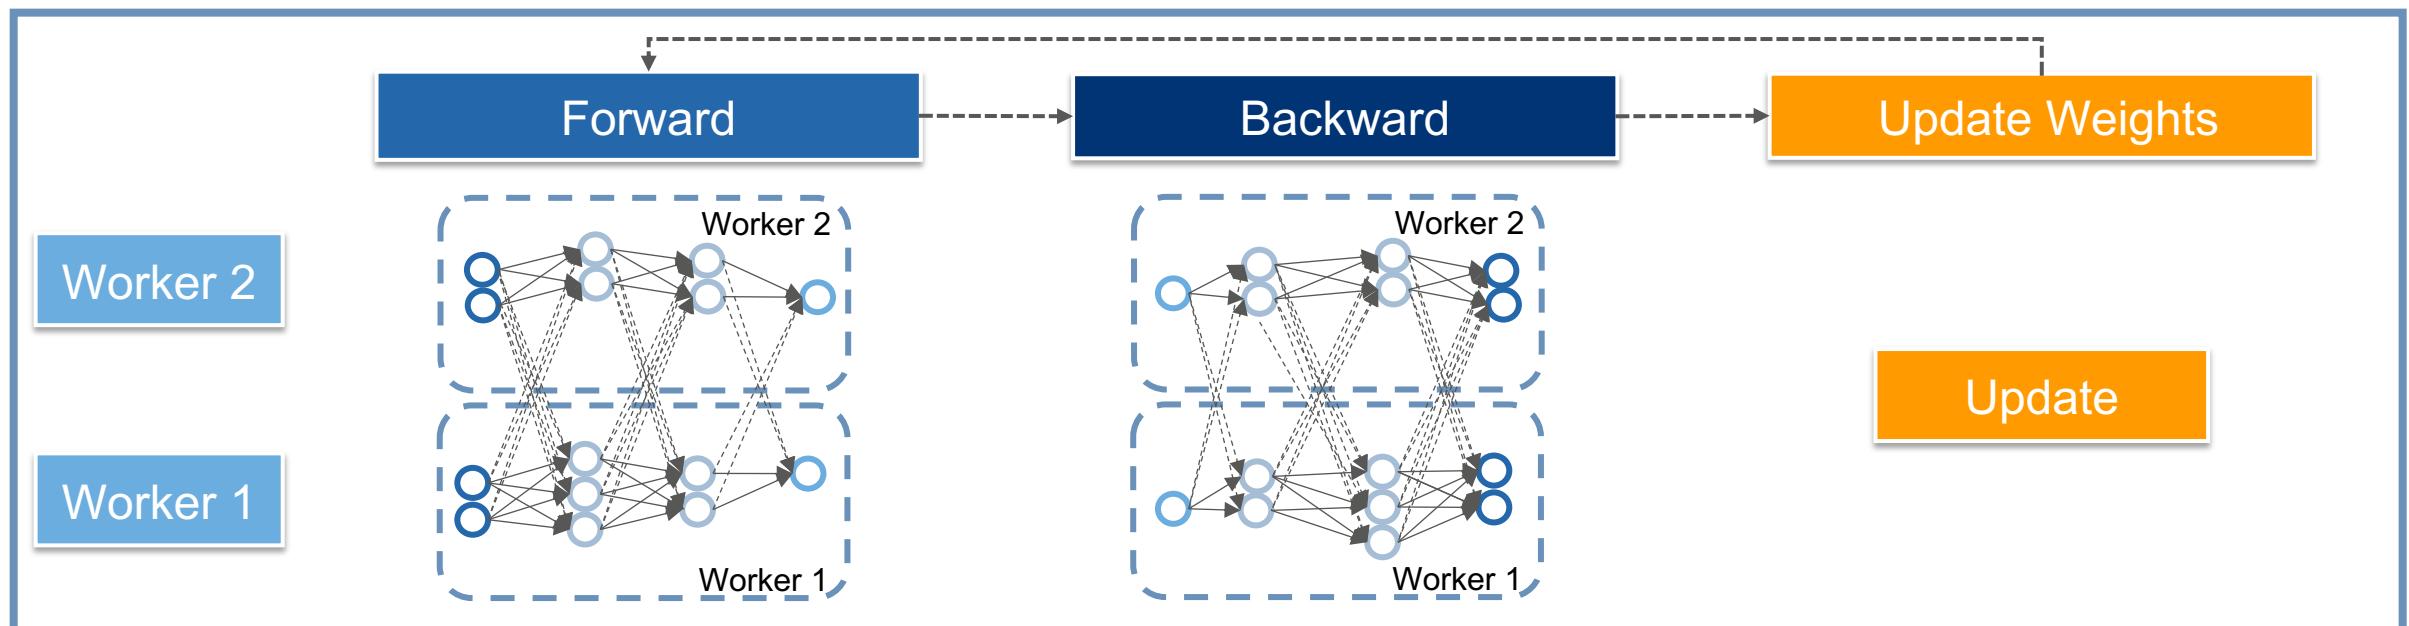
Output:

```
Total parameters in model: 1,444,261,998
end of epoch 1 | time: 28.42s | valid loss 1.22
end of epoch 2 | time: 27.36s | valid loss 0.25
end of epoch 3 | time: 27.49s | valid loss 0.23
End of training | test loss 0.20
```

### 3. Distributed Training of Neural Networks

## Tensor Parallel Training of Neural Networks

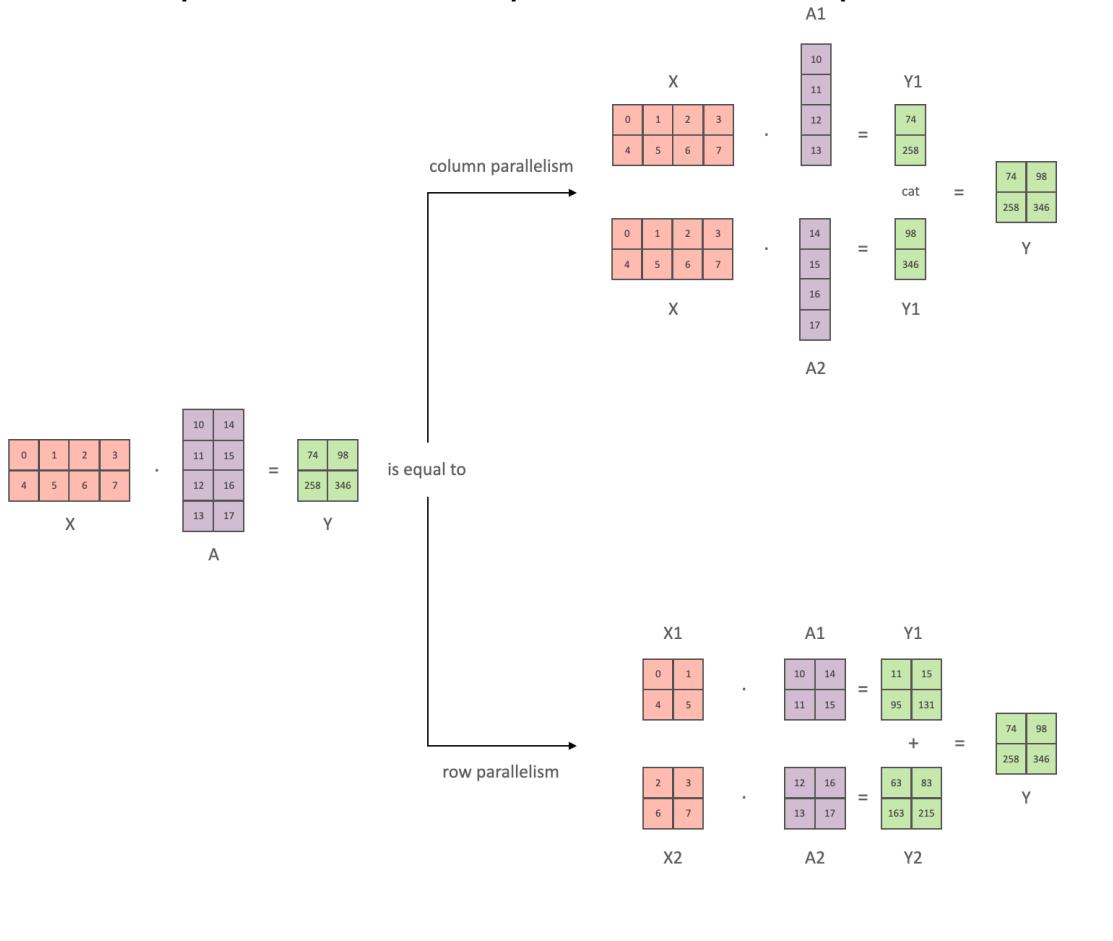
- Each operator or layer of the neural network split between the workers.
- There is a single copy of weight parameters shared between the workers.
- Goal: Make feasible models that are larger than the memory capacity of a single GPU.
- **Collective (all-to-all)** communication of partial activations and gradients of each worker.



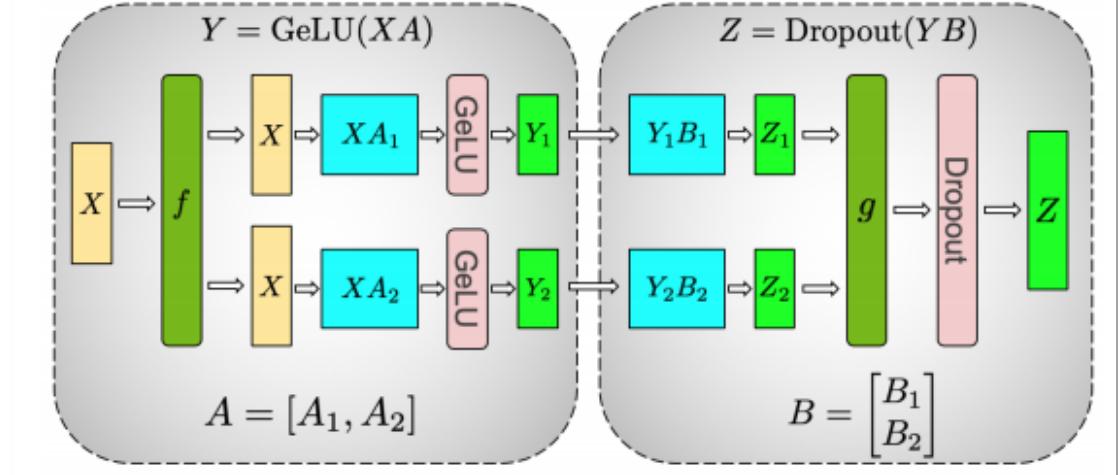
### 3. Distributed Training of Neural Networks

## Tensor Parallel Training of Neural Networks

Matrix multiplication can be split between multiple GPUs



Source: <https://huggingface.co/docs/transformers/parallelism#naive-model-parallelism-vertical-and-pipeline-parallelism>



(a) MLP

Source: Narayanan et al., 2021, arXiv:2104.04473

## Combination of Data, Pipeline and Tensor Parallelism – Megatron-LM

### How do we navigate this configuration space?

Degree of pipeline, tensor,  
and data parallelism  
  
Pipelining schedule  
  
Global batch size  
  
Microbatch size

Each of these choices influence  
amount of communication, size of  
pipeline bubble, memory footprint

Tensor, pipeline, and data parallelism  
can be composed to scale with a trillion  
parameters to thousands of GPUs.

Source: Narayanan et al., 2021, arXiv:2104.04473  
<https://github.com/NVIDIA/Megatron-LM>  
[https://www.youtube.com/watch?v=MO\\_CESBOXgo](https://www.youtube.com/watch?v=MO_CESBOXgo)

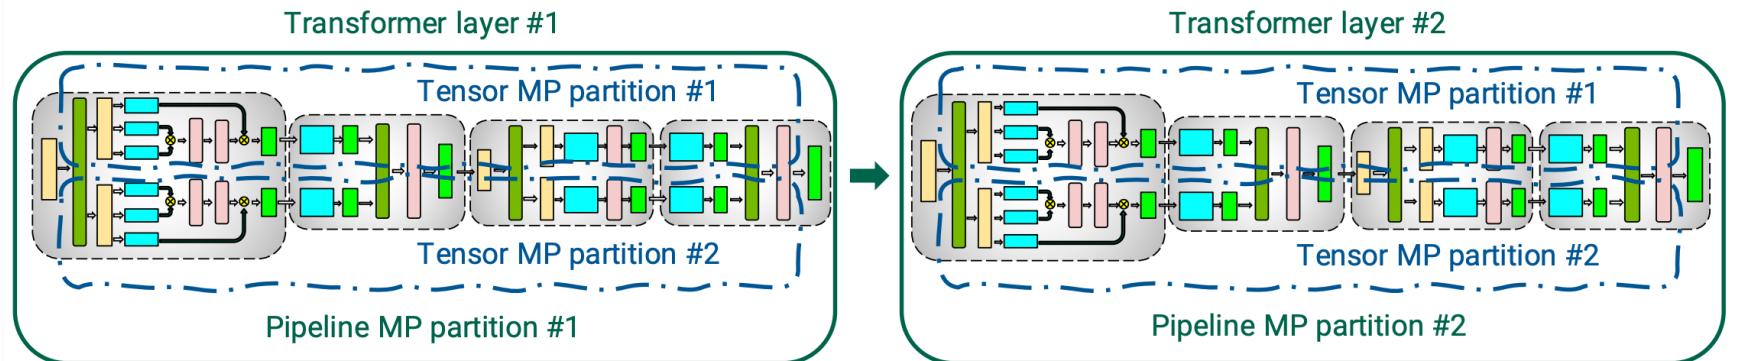


Figure 2: Combination of tensor and pipeline model parallelism (MP) used in this work for transformer-based models.

- Pytorch tutorials on distributed training of transformer models:
  - Data and pipeline parallelism: [https://pytorch.org/tutorials/advanced/ddp\\_pipeline.html](https://pytorch.org/tutorials/advanced/ddp_pipeline.html)
- Mesh TensorFlow data and model parallelism:
  - GitHub: <https://github.com/tensorflow/mesh>
  - Paper: Shazeer et al., 2018, <https://arxiv.org/pdf/1811.02084.pdf>
- DeepSpeed (PyTorch deep learning optimization library by Microsoft):
  - Pipeline parallelism: <https://www.deepspeed.ai/tutorials/pipeline/>
  - ZeRO data parallelism: <https://deepspeed.readthedocs.io/en/latest/zero3.html>
- FairScale (PyTorch library for high performance large scale training by Facebook Research):
  - Pipeline parallelism: <https://fairscale.readthedocs.io/en/latest/tutorials/pipe.html>
  - ZeRO data parallelism: <https://github.com/facebookresearch/fairscale/#optimizer-state-sharding-zero>

## 1. Introduction to the LRZ AI Systems

---

- ✓ Overview of the LRZ AI Systems
- ✓ Access to the LRZ AI Systems
- ✓ NVIDIA NGC Cloud
- ✓ Introduction to Enroot Containers
- ✓ Interactive and Batch Jobs
- ✓ Open on Demand
- ✓ Exercise: Run a job with an Enroot container

## 2. Fundamentals of Deep Learning

---

- ✓ Introduction to Neural Networks
- ✓ Introduction to Convolutional Neural Networks
- ✓ Exercises: Train CNNs on a GPU
- ✓ Introduction to Transformers
- ✓ Exercise: Train a Transformer on a GPU

## 3. Distributed Training of Neural Networks

---

- ✓ Data Parallel Training
- ✓ Exercise: DP Training of CNN on 2 GPUs
- ✓ Model Parallel Training: Pipeline Parallel and Tensor Parallel
- ✓ Exercise: PP Training of Transformer on 2 GPUs

Please visit

<https://survey.lrz.de/index.php/538854?lang=en>

and rate this course.

Your feedback is highly appreciated!

Thank you!



# ANNOUNCEMENT



<https://app1.edoobox.com/en/LRZ/>

| All offers  | Search term             | Search                | English | Login |
|---|-------------------------|-----------------------|---------|-------|
| <a href="#">Introduction to ANSYS CFX on LRZ HPC Systems</a>                                      | 16.03.2023 – 27.04.2023 | ONLINE                | 0.00 €  | 62    |
| <a href="#">AI Training Series - Orientation Session</a>  | 12.04.2023 – 13.04.2023 | HYBRID: ONLINE/LRZ    | 0.00 €  | 30    |
| <a href="#">EuroCC AI for Science Bootcamp (apply via openhackathons.org)</a>                     | 17.04.2023 – 18.04.2023 | ONLINE                | 0.00 €  | 100   |
| <a href="#">Introduction to LRZ HPC Systems with Focus on CFD Workflows</a>                       | 19.04.2023 – 19.04.2023 | ONLINE                | 0.00 €  | 51    |
| <a href="#">AI Training Series - Intro to Container Technology &amp; Application to AI at LRZ</a> | 26.04.2023 – 26.04.2023 | HYBRID: ONLINE/LRZ    | 0.00 €  | 43    |
| <a href="#">AI Training Series - Introduction to the LRZ AI Systems</a>                           | 03.05.2023 – 03.05.2023 | ONLINE                | 0.00 €  | 47    |
| <a href="#">AI Training Series - Intel AI Workshop - Accelerated Machine Learning with Intel</a>  | 04.05.2023 – 04.05.2023 | ONLINE                | 0.00 €  | 88    |
| <a href="#">SuperMUC-NG Status and Results Workshop</a>   | 09.05.2023 – 11.05.2023 | ONLINE                | 0.00 €  | 192   |
| <a href="#">AI Training Series - Introduction to the LRZ Linux Cluster</a>                        | 09.05.2023 – 09.05.2023 | HYBRID: ONLINE/LRZ    | 0.00 €  | 73    |
| <a href="#">Working with Noisy Quantum Computers</a>  | 11.05.2023 – 11.05.2023 | Leibniz Rechenzentrum | 0.00 €  | 3     |
| <a href="#">AI Training Series - High Performance Data Analytics Using R at LRZ</a>               | 12.05.2023 – 12.05.2023 | HYBRID: ONLINE/LRZ    | 0.00 €  | 83    |
| <a href="#">EuroCC N-Ways to GPU Programming Bootcamp (apply via openhackathons.org)</a>          | 15.05.2023 – 16.05.2023 | ONLINE                | 0.00 €  | 100   |
| <a href="#">AI Training Series - Intel AI Workshop - Accelerated Deep Learning with Intel</a>     | 16.05.2023 – 16.05.2023 | HYBRID: ONLINE/LRZ    | 0.00 €  | 95    |
| <a href="#">AI Training Series - Introduction to the LRZ Compute Cloud</a>                        | 17.05.2023 – 17.05.2023 | HYBRID: ONLINE/LRZ    | 0.00 €  | 85    |

Thank you for your  
attention! 😊



Maja Piskac  
Leibniz Supercomputing Centre  
LRZ AI & Big Data Team  
[maja.piskac@lrz.de](mailto:maja.piskac@lrz.de)